

Africast-Time Series Analysis & Forecasting Using R

<https://workshop.f4sg.org/africast/>

1. Basics of time series and data structures



Outline


- 1 Introduction to forecasting
- 2 Time series data
- 3 tsibble objects
- 4 Example: Australian prison population
- 5 Example: Australian pharmaceutical sales

Outline

- 1 Introduction to forecasting
- 2 Time series data
- 3 tsibble objects
- 4 Example: Australian prison population
- 5 Example: Australian pharmaceutical sales

What is a forecast?

Estimating how a sequence of observations will continue into the future, using all information available now.



What data do we need?

- 1 Historical data on the variable to forecast
- 2 Past and future values of predictors
 - ▶ Deterministic (e.g. holidays)
 - ▶ Stochastic (e.g. temperature)
- 3 Expert knowledge and context
 - ▶ New information

Why forecast?

Why do you use forecasts?


Why forecast?

Why do you use forecasts?

Decision	Forecast
Build a new hospital in the next 10 years?	?
Call centre staff needed next week?	?
Vaccine doses needed next month?	?

- Supports planning
- Informs decisions
- Provides evidence

Forecasting workflow



Packages for this course

tidyverts.org



Packages for this course

```
# Data, graphics and forecasting
```

```
library(fpp3)
```

```
-- Attaching packages ----- fpp3 1.0.3 --
```

v tibble	3.3.1	v tsibble	1.2.0.9000
v dplyr	1.2.1	v tsibbledata	0.4.1
v tidyr	1.3.2	v ggtime	0.2.0.9000
v lubridate	1.9.5	v feasts	0.5.0
v ggplot2	4.0.3	v fable	0.5.0.9000

Packages for this course

```
# Data, graphics and forecasting
```

```
library(fpp3)
```

```
-- Attaching packages ----- fpp3 1.0.3 --
```

v tibble	3.3.1	v tsibble	1.2.0.9000
v dplyr	1.2.1	v tsibbledata	0.4.1
v tidyr	1.3.2	v ggtime	0.2.0.9000
v lubridate	1.9.5	v feasts	0.5.0
v ggplot2	4.0.3	v fable	0.5.0.9000

```
# Time classes for the index (load after fpp3)
```

```
library(mixtime)
```

Outline

- 1 Introduction to forecasting
- 2 Time series data
- 3 tsibble objects
- 4 Example: Australian prison population
- 5 Example: Australian pharmaceutical sales

Time series data

- Four-yearly Olympic winning times
- Annual Google profits
- Quarterly Australian beer production
- Monthly rainfall
- Weekly retail sales
- Daily IBM stock prices
- Hourly electricity demand
- 5-minute freeway traffic counts
- Time-stamped stock transaction data

Your time series

What time series do you work with?

What time series do you work with?

- What is measured?
- How often?
- One series, or many?

What makes it a time series?

Observations + **when** they happened.

- The time column is the **index**
- Observations are ordered in time
- The frequency matters

Outline

- 1 Introduction to forecasting
- 2 Time series data
- 3 tsibble objects
- 4 Example: Australian prison population
- 5 Example: Australian pharmaceutical sales

A `tsibble` has three parts

A **tsibble** stores one or more time series in a table.

- **Index:** when
- **Key:** which series (optional)
- **Measures:** what was observed

Works with tidyverse functions.

tsibble objects

```
global_economy
```

```
# A tsibble: 15,150 x 6 [1Y]
```

```
# Key:      Country [263]
```

	Year	Country	GDP	Imports	Exports	Population
	<dbl>	<fct>	<dbl>	<dbl>	<dbl>	<dbl>
1	1960	Afghanistan	537777811.	7.02	4.13	8996351
2	1961	Afghanistan	548888896.	8.10	4.45	9166764
3	1962	Afghanistan	546666678.	9.35	4.88	9345868
4	1963	Afghanistan	751111191.	16.9	9.17	9533954
5	1964	Afghanistan	800000044.	18.1	8.89	9731361
6	1965	Afghanistan	1006666638.	21.4	11.3	9938414
7	1966	Afghanistan	1399999967.	18.6	8.57	10152331
8	1967	Afghanistan	1673333418.	14.2	6.77	10372630
9	1968	Afghanistan	1373333367.	15.2	8.90	10604346
10	1969	Afghanistan	1408888922.	15.0	10.1	10854428

```
# i 15,140 more rows
```

tsibble objects

```
global_economy
```

```
# A tsibble: 15,150 x 6 [1Y]
```

```
# Key:           Country [263]
```

	Year	Country	GDP	Imports	Exports	Population
	Index	<fct>	<dbl>	<dbl>	<dbl>	<dbl>
1	1960	Afghanistan	537777811.	7.02	4.13	8996351
2	1961	Afghanistan	548888896.	8.10	4.45	9166764
3	1962	Afghanistan	546666678.	9.35	4.88	9345868
4	1963	Afghanistan	751111191.	16.9	9.17	9533954
5	1964	Afghanistan	800000044.	18.1	8.89	9731361
6	1965	Afghanistan	1006666638.	21.4	11.3	9938414
7	1966	Afghanistan	1399999967.	18.6	8.57	10152331
8	1967	Afghanistan	1673333418.	14.2	6.77	10372630
9	1968	Afghanistan	1373333367.	15.2	8.90	10604346
10	1969	Afghanistan	1408888922.	15.0	10.1	10854428

```
# i 15,140 more rows
```

tsibble objects

```
global_economy
```

```
# A tsibble: 15,150 x 6 [1Y]
```

```
# Key:      Country [263]
```

	Year	Country	GDP	Imports	Exports	Population
	Index	Key	<dbl>	<dbl>	<dbl>	<dbl>
1	1960	Afghanistan	5377777811.	7.02	4.13	8996351
2	1961	Afghanistan	5488888896.	8.10	4.45	9166764
3	1962	Afghanistan	5466666678.	9.35	4.88	9345868
4	1963	Afghanistan	7511111191.	16.9	9.17	9533954
5	1964	Afghanistan	8000000044.	18.1	8.89	9731361
6	1965	Afghanistan	10066666638.	21.4	11.3	9938414
7	1966	Afghanistan	13999999967.	18.6	8.57	10152331
8	1967	Afghanistan	16733333418.	14.2	6.77	10372630
9	1968	Afghanistan	13733333367.	15.2	8.90	10604346
10	1969	Afghanistan	14088888922.	15.0	10.1	10854428

```
# i 15,140 more rows
```

tsibble objects

```
global_economy
```

```
# A tsibble: 15,150 x 6 [1Y]
```

```
# Key:      Country [263]
```

	Year	Country	GDP	Imports	Exports	Population
Index	Key	Measures				
1	1960	Afghanistan	537777811.	7.02	4.13	8996351
2	1961	Afghanistan	548888896.	8.10	4.45	9166764
3	1962	Afghanistan	546666678.	9.35	4.88	9345868
4	1963	Afghanistan	751111191.	16.9	9.17	9533954
5	1964	Afghanistan	800000044.	18.1	8.89	9731361
6	1965	Afghanistan	1006666638.	21.4	11.3	9938414
7	1966	Afghanistan	1399999967.	18.6	8.57	10152331
8	1967	Afghanistan	1673333418.	14.2	6.77	10372630
9	1968	Afghanistan	1373333367.	15.2	8.90	10604346
10	1969	Afghanistan	1408888922.	15.0	10.1	10854428

```
# i 15,140 more rows
```

tsibble objects

tourism

```
# A tsibble: 24,320 x 5 [1Q]
# Key:           Region, State, Purpose [304]
  Quarter Region State Purpose Trips
  <qtr> <chr>    <chr> <chr>    <dbl>
1 1998 Q1 Adelaide SA      Business 135.
2 1998 Q2 Adelaide SA      Business 110.
3 1998 Q3 Adelaide SA      Business 166.
4 1998 Q4 Adelaide SA      Business 127.
5 1999 Q1 Adelaide SA      Business 137.
6 1999 Q2 Adelaide SA      Business 200.
7 1999 Q3 Adelaide SA      Business 169.
8 1999 Q4 Adelaide SA      Business 134.
9 2000 Q1 Adelaide SA      Business 154.
10 2000 Q2 Adelaide SA      Business 169.
# i 24,310 more rows
```

Domestic visitor
nights
(thousands).

tsibble objects

tourism

```
# A tsibble: 24,320 x 5 [1Q]
# Key:           Region, State, Purpose [304]
   Quarter Region State Purpose Trips
   <chr>    <chr>  <chr> <chr>    <dbl>
1 1998 Q1 Adelaide SA      Business 135.
2 1998 Q2 Adelaide SA      Business 110.
3 1998 Q3 Adelaide SA      Business 166.
4 1998 Q4 Adelaide SA      Business 127.
5 1999 Q1 Adelaide SA      Business 137.
6 1999 Q2 Adelaide SA      Business 200.
7 1999 Q3 Adelaide SA      Business 169.
8 1999 Q4 Adelaide SA      Business 134.
9 2000 Q1 Adelaide SA      Business 154.
10 2000 Q2 Adelaide SA      Business 169.
# i 24,310 more rows
```

Domestic visitor
nights
(thousands).

tsibble objects

tourism

```
# A tsibble: 24,320 x 5 [1Q]
# Key:      Region, State, Purpose [304]
  Quarter Region State Purpose Trips
  <dbl>   <chr>   <chr>   <chr>   <dbl>
1 1998 Q1 Adelaide SA      Business 135.
2 1998 Q2 Adelaide SA      Business 110.
3 1998 Q3 Adelaide SA      Business 166.
4 1998 Q4 Adelaide SA      Business 127.
5 1999 Q1 Adelaide SA      Business 137.
6 1999 Q2 Adelaide SA      Business 200.
7 1999 Q3 Adelaide SA      Business 169.
8 1999 Q4 Adelaide SA      Business 134.
9 2000 Q1 Adelaide SA      Business 154.
10 2000 Q2 Adelaide SA      Business 169.
# i 24,310 more rows
```

Domestic visitor
nights
(thousands).

tsibble objects

tourism

```
# A tsibble: 24,320 x 5 [1Q]
```

```
# Key:      Region, State, Purpose [304]
```

```
  Quarter Region State Purpose Trips
```

```
  Index  Keys Measure
```

```
1 1998 Q1 Adelaide SA Business 135.
2 1998 Q2 Adelaide SA Business 110.
3 1998 Q3 Adelaide SA Business 166.
4 1998 Q4 Adelaide SA Business 127.
5 1999 Q1 Adelaide SA Business 137.
6 1999 Q2 Adelaide SA Business 200.
7 1999 Q3 Adelaide SA Business 169.
8 1999 Q4 Adelaide SA Business 134.
9 2000 Q1 Adelaide SA Business 154.
10 2000 Q2 Adelaide SA Business 169.
```

```
# i 24,310 more rows
```

Domestic visitor
nights
(thousands).

The index needs a time class

z

```
# A tibble: 5 x 2
  Month      Observation
  <chr>          <dbl>
1 2019 Jan           50
2 2019 Feb           23
3 2019 Mar           34
4 2019 Apr           30
5 2019 May           25
```

The index needs a time class

z

```
# A tibble: 5 x 2
  Month      Observation
  <chr>          <dbl>
1 2019 Jan           50
2 2019 Feb           23
3 2019 Mar           34
4 2019 Apr           30
5 2019 May           25
```

- "2019 Jan" is just text
- A time class knows the frequency

Time classes with `mixtime`

Three kinds of time:

Kind	Meaning	Function	Example
Linear	Point on a timeline	<code>linear_time()</code>	2026 Oct
Cyclical	Position in a cycle	<code>cyclical_time()</code>	Mon
Duration	Length of time	<code>duration()</code>	3 months



Time classes with `mixtime`

Three kinds of time:

Kind	Meaning	Function	Example
Linear	Point on a timeline	<code>linear_time()</code>	2026 Oct
Cyclical	Position in a cycle	<code>cyclical_time()</code>	Mon
Duration	Length of time	<code>duration()</code>	3 months

The time index for a tsibble is typically **linear time**.



Linear time

Common frequencies have convenient functions:

Frequency	Function
Annual	<code>year()</code>
Quarterly	<code>yearquarter()</code>
Monthly	<code>yearmonth()</code>
Weekly	<code>yearweek()</code>
Daily	<code>date()</code>
Seconds	<code>datetime()</code>

Linear time

Often data is provided in a date format, even if it isn't daily.

```
d <- as.Date("2026-10-01")
yearquarter(d)
## <mixtime[1]>
## [1] 2026 Q4
yearmonth(d) + 0:2
## <mixtime[3]>
## [1] 2026 Oct 2026 Nov 2026 Dec
```

Linear time

Often data is provided in a date format, even if it isn't daily.

```
d <- as.Date("2026-10-01")
yearquarter(d)
## <mixtime[1]>
## [1] 2026 Q4
yearmonth(d) + 0:2
## <mixtime[3]>
## [1] 2026 Oct 2026 Nov 2026 Dec
```

Multiple granularities can share one vector:

```
c(yearquarter(d), yearmonth(d), date(d))
## <mixtime[3]>
## [1] 2026 Q4      2026 Oct      2026-10-05
```

Making a tsibble

```
z |>  
  mutate(Month = yearmonth(Month)) |>  
  as_tsibble(index = Month)
```

```
# A tsibble: 5 x 2 [1M]  
  Month      Observation  
  <mixtime>      <dbl>  
1 2019 Jan           50  
2 2019 Feb           23  
3 2019 Mar           34  
4 2019 Apr           30  
5 2019 May           25
```

Custom granularity

Any time step with `linear_time()`:

```
x <- datetime("2024-01-01 00:00:00") + minutes(0:5) * 15L
linear_time(x, chronon = cal_gregorian$minute(15L))
## <mixtime[6]>
## [1] 2024-01-01 00:00 2024-01-01 00:15 2024-01-01 00:30 2024-01-01 00:45 2024-01-01 01:00
## [6] 2024-01-01 01:15
```

Custom granularity

Any time step with `linear_time()`:

```
x <- datetime("2024-01-01 00:00:00") + minutes(0:5) * 15L
linear_time(x, chronon = cal_gregorian$minute(15L))
## <mixtime[6]>
## [1] 2024-01-01 00:00 2024-01-01 00:15 2024-01-01 00:30 2024-01-01 00:45 2024-01-01 01:00
## [6] 2024-01-01 01:15
```

- `chronon`: the size of one time step
- `cal_gregorian$minute(15L)`: 15 minutes
- `cal_gregorian$hour(6L)`: 6 hours

Timezones

```
t <- as.POSIXct("2026-10-05 22:30:00", tz = "UTC")
datetime(t, tz = "Africa/Lagos")
## <mixtime[1]>
## [1] 2026-10-05 23:30:00 WAT
datetime(t, tz = "Africa/Nairobi")
## <mixtime[1]>
## [1] 2026-10-06 01:30:00 EAT
date(t, tz = "Africa/Nairobi")
## <mixtime[1]>
## [1] 2026-10-06 EAT
```

Timezones

```
t <- as.POSIXct("2026-10-05 22:30:00", tz = "UTC")
datetime(t, tz = "Africa/Lagos")
## <mixtime[1]>
## [1] 2026-10-05 23:30:00 WAT
datetime(t, tz = "Africa/Nairobi")
## <mixtime[1]>
## [1] 2026-10-06 01:30:00 EAT
date(t, tz = "Africa/Nairobi")
## <mixtime[1]>
## [1] 2026-10-06 EAT
```

- Same moment, different local time zone.
- The date can change!

Cyclical time

Cyclical time describes positions within a cycle:

```
month_of_year(d)
## <mixtime[1]>
## [1] Oct
day_of_week(d)
## <mixtime[1]>
## [1] Mon
time_of_day(t)
## <mixtime[1]>
## [1] 22:30:00
```

Your turn

What day of the week
were you born on?

Cyclical time

Cyclical time describes positions within a cycle:

```
month_of_year(d)
## <mixtime[1]>
## [1] Oct
day_of_week(d)
## <mixtime[1]>
## [1] Mon
time_of_day(t)
## <mixtime[1]>
## [1] 22:30:00
```

Your turn

What day of the week
were you born on?

- Also `week_of_year()`, `day_of_month()`, `day_of_year()`
- Useful for seasonal patterns

Cyclical time

Cyclical time describes positions within a cycle:

```
month_of_year(d)
## <mixtime[1]>
## [1] Oct
day_of_week(d)
## <mixtime[1]>
## [1] Mon
time_of_day(t)
## <mixtime[1]>
## [1] 22:30:00
```

Your turn

What day of the week
were you born on?

- Also `week_of_year()`, `day_of_month()`, `day_of_year()`
- Useful for seasonal patterns

Durations

Durations describe a length of time:

```
months(3)
## <mixtime[1]>
## [1] 3.0 months
```

- Also `years()`, `quarters()`, `weeks()`, `days()`, `hours()`, `minutes()`, `seconds()`

Durations

Durations describe a length of time:

```
months(3)
## <mixtime[1]>
## [1] 3.0 months
```

- Also `years()`, `quarters()`, `weeks()`, `days()`, `hours()`, `minutes()`, `seconds()`

They are useful for arithmetic.

```
yearmonth(d) + months(0:2)
## <mixtime[3]>
## [1] 2026 Oct 2026 Nov 2026 Dec
yearmonth("2026 Oct") - yearmonth("2025 Jan")
## <mixtime[1]>
## [1] 21.0 months
```

Durations

Durations describe a length of time:

```
months(3)
## <mixtime[1]>
## [1] 3.0 months
```

- Also `years()`, `quarters()`, `weeks()`, `days()`, `hours()`, `minutes()`, `seconds()`

They are useful for arithmetic.

```
yearmonth(d) + months(0:2)
## <mixtime[3]>
## [1] 2026 Oct 2026 Nov 2026 Dec
yearmonth("2026 Oct") - yearmonth("2025 Jan")
## <mixtime[1]>
## [1] 21.0 months
```

Outline

- 1 Introduction to forecasting
- 2 Time series data
- 3 tsibble objects
- 4 Example: Australian prison population
- 5 Example: Australian pharmaceutical sales

Australian prison population



Read a csv file

```
prison <- readr::read_csv("data/prison_population.csv")
```

```
# A tibble: 3,072 x 6
```

	date	state	gender	legal	indigenous	count
	<date>	<chr>	<chr>	<chr>	<chr>	<dbl>
1	2005-03-01	ACT	Female	Remanded	ATSI	0
2	2005-03-01	ACT	Female	Remanded	Other	2
3	2005-03-01	ACT	Female	Sentenced	ATSI	0
4	2005-03-01	ACT	Female	Sentenced	Other	0
5	2005-03-01	ACT	Male	Remanded	ATSI	7
6	2005-03-01	ACT	Male	Remanded	Other	58
7	2005-03-01	ACT	Male	Sentenced	ATSI	0
8	2005-03-01	ACT	Male	Sentenced	Other	0
9	2005-03-01	NSW	Female	Remanded	ATSI	51
10	2005-03-01	NSW	Female	Remanded	Other	131

```
# i 3,062 more rows
```

Convert to a time class

```
prison <- readr::read_csv("data/prison_population.csv") |>  
  mutate(quarter = yearquarter(date))
```

A tibble: 3,072 x 7

	date	state	gender	legal	indigenous	count	quarter
	<date>	<chr>	<chr>	<chr>	<chr>	<dbl>	<mixtime>
1	2005-03-01	ACT	Female	Remanded	ATSI	0	2005 Q1
2	2005-03-01	ACT	Female	Remanded	Other	2	2005 Q1
3	2005-03-01	ACT	Female	Sentenced	ATSI	0	2005 Q1
4	2005-03-01	ACT	Female	Sentenced	Other	0	2005 Q1
5	2005-03-01	ACT	Male	Remanded	ATSI	7	2005 Q1
6	2005-03-01	ACT	Male	Remanded	Other	58	2005 Q1
7	2005-03-01	ACT	Male	Sentenced	ATSI	0	2005 Q1
8	2005-03-01	ACT	Male	Sentenced	Other	0	2005 Q1
9	2005-03-01	NSW	Female	Remanded	ATSI	51	2005 Q1
10	2005-03-01	NSW	Female	Remanded	Other	131	2005 Q1

i 3,062 more rows

Remove the old date column

```
prison <- readr::read_csv("data/prison_population.csv") |>  
  mutate(quarter = yearquarter(date)) |>  
  select(-date)
```

A tibble: 3,072 x 6

	state	gender	legal	indigenous	count	quarter
	<chr>	<chr>	<chr>	<chr>	<dbl>	<mixtime>
1	ACT	Female	Remanded	ATSI	0	2005 Q1
2	ACT	Female	Remanded	Other	2	2005 Q1
3	ACT	Female	Sentenced	ATSI	0	2005 Q1
4	ACT	Female	Sentenced	Other	0	2005 Q1
5	ACT	Male	Remanded	ATSI	7	2005 Q1
6	ACT	Male	Remanded	Other	58	2005 Q1
7	ACT	Male	Sentenced	ATSI	0	2005 Q1
8	ACT	Male	Sentenced	Other	0	2005 Q1
9	NSW	Female	Remanded	ATSI	51	2005 Q1
10	NSW	Female	Remanded	Other	131	2005 Q1

i 3,062 more rows

Declare the index and keys

```
prison <- readr::read_csv("data/prison_population.csv") |>
  mutate(quarter = yearquarter(date)) |>
  select(-date) |>
  as_tsibble(
    index = quarter,
    key = c(state, gender, legal, indigenous)
  )
```

```
# A tsibble: 3,072 x 6 [1Q]
```

```
# Key:      state, gender, legal, indigenous [64]
```

	state	gender	legal	indigenous	count	quarter
	<chr>	<chr>	<chr>	<chr>	<dbl>	<mixtime>
1	ACT	Female	Remanded	ATSI	0	2005 Q1
2	ACT	Female	Remanded	ATSI	1	2005 Q2
3	ACT	Female	Remanded	ATSI	0	2005 Q3
4	ACT	Female	Remanded	ATSI	0	2005 Q4
5	ACT	Female	Remanded	ATSI	1	2006 Q1
6	ACT	Female	Remanded	ATSI	1	2006 Q2

Outline

- 1 Introduction to forecasting
- 2 Time series data
- 3 tsibble objects
- 4 Example: Australian prison population
- 5 Example: Australian pharmaceutical sales

Australian Pharmaceutical Benefits Scheme



Australian Pharmaceutical Benefits Scheme

The **PBS** is the Australian government drug subsidy scheme.

Australian Pharmaceutical Benefits Scheme

The **PBS** is the Australian government drug subsidy scheme.

- Subsidises many pharmacy drugs
- Costs nearly 1% of GDP
- Budget relies on forecasts of drug usage
- Split by drug type (ATC2 x84), concession (x2) and patient type (x2)

Working with tsibble objects

PBS

```
# A tsibble: 67,596 x 9 [1M]
# Key:      Concession, Type, ATC1, ATC2 [336]
   Month Concession  Type  ATC1 ATC1_desc ATC2 ATC2_desc Scripts  Cost
   <mth> <chr>      <chr> <chr> <chr>    <chr> <chr>    <dbl> <dbl>
1 1991 Jul  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 18228 67877
2 1991 Aug  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 15327 57011
3 1991 Sep  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 14775 55020
4 1991 Oct  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 15380 57222
5 1991 Nov  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 14371 52120
6 1991 Dec  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 15028 54299
7 1992 Jan  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 11040 39753
8 1992 Feb  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 15165 54405
9 1992 Mar  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 16898 61108
10 1992 Apr  Concessional Co-pay~ A    Alimenta~ A01  STOMATOL~ 18141 65356
# i 67,586 more rows
```

filter(): choose rows

```
PBS |>  
  filter(ATC2 == "A10")
```

```
# A tibble: 816 x 9 [1M]
```

```
# Key:      Concession, Type, ATC1, ATC2 [4]
```

	Month	Concession	Type	ATC1	ATC1_desc	ATC2	ATC2_desc	Scripts	Cost
	<mth>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>
1	1991 Jul	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	89733	2.09e6
2	1991 Aug	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	77101	1.80e6
3	1991 Sep	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	76255	1.78e6
4	1991 Oct	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	78681	1.85e6
5	1991 Nov	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	70554	1.69e6
6	1991 Dec	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	75814	1.84e6
7	1992 Jan	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	64186	1.56e6
8	1992 Feb	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	75899	1.73e6
9	1992 Mar	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	89445	2.05e6
10	1992 Apr	Concessional	Co-pa~	A	Alimenta~	A10	ANTIDIAB~	97315	2.23e6

```
# i 806 more rows
```

select(): choose columns

```
PBS |>
  filter(ATC2 == "A10") |>
  select(Month, Concession, Type, Cost)
```

```
# A tibble: 816 x 4 [1M]
```

```
# Key:      Concession, Type [4]
```

	Month	Concession	Type	Cost
	<mth>	<chr>	<chr>	<dbl>
1	1991 Jul	Concessional	Co-payments	2092878
2	1991 Aug	Concessional	Co-payments	1795733
3	1991 Sep	Concessional	Co-payments	1777231
4	1991 Oct	Concessional	Co-payments	1848507
5	1991 Nov	Concessional	Co-payments	1686458
6	1991 Dec	Concessional	Co-payments	1843079
7	1992 Jan	Concessional	Co-payments	1564702
8	1992 Feb	Concessional	Co-payments	1732508
9	1992 Mar	Concessional	Co-payments	2046102
10	1992 Apr	Concessional	Co-payments	2225977

```
# i 806 more rows
```

summarise(): combine over keys

```
PBS |>
  filter(ATC2 == "A10") |>
  select(Month, Concession, Type, Cost) |>
  summarise(total_cost = sum(Cost))
```

```
# A tsibble: 204 x 2 [1M]
```

	Month	total_cost
	<mth>	<dbl>
1	1991 Jul	3526591
2	1991 Aug	3180891
3	1991 Sep	3252221
4	1991 Oct	3611003
5	1991 Nov	3565869
6	1991 Dec	4306371
7	1992 Jan	5088335
8	1992 Feb	2814520
9	1992 Mar	2985811
10	1992 Apr	3204780

```
# i 194 more rows
```

mutate(): create variables

```
PBS |>  
  filter(ATC2 == "A10") |>  
  select(Month, Concession, Type, Cost) |>  
  summarise(total_cost = sum(Cost)) |>  
  mutate(total_cost = total_cost / 1e6)
```

```
# A tsibble: 204 x 2 [1M]
```

	Month	total_cost
	<mth>	<dbl>
1	1991 Jul	3.53
2	1991 Aug	3.18
3	1991 Sep	3.25
4	1991 Oct	3.61
5	1991 Nov	3.57
6	1991 Dec	4.31
7	1992 Jan	5.09
8	1992 Feb	2.81
9	1992 Mar	2.99
10	1992 Apr	3.20

Save the result

```
PBS |>
  filter(ATC2 == "A10") |>
  select(Month, Concession, Type, Cost) |>
  summarise(total_cost = sum(Cost)) |>
  mutate(total_cost = total_cost / 1e6) -> a10
```

```
# A tsibble: 204 x 2 [1M]
```

	Month	total_cost
	<mth>	<dbl>
1	1991 Jul	3.53
2	1991 Aug	3.18
3	1991 Sep	3.25
4	1991 Oct	3.61
5	1991 Nov	3.57
6	1991 Dec	4.31
7	1992 Jan	5.09
8	1992 Feb	2.81
9	1992 Mar	2.99
10	1992 Apr	3.20

Your turn

Using tourism

- 1 Find the quarterly total `Trips` for each State
- 2 **Challenge:** Which State had the most `Trips` over the whole period?

Using PBS

- 3 Find the total monthly `Scripts` for ATC2 "H02"

Your turn

Using tourism

- 1 Find the quarterly total `Trips` for each State
- 2 **Challenge:** Which State had the most `Trips` over the whole period?

Using PBS

- 3 Find the total monthly `Scripts` for ATC2 "H02"

Hints

- `summarise()` always keeps the index
- `group_by()` the keys you want to keep
- `as_tibble()` drops the index, to summarise over time

Recap

- A tsibble has an **index**, **keys** and **measures**
- The index needs a time class (`mixtime`)
- Time can be linear, cyclical or a duration
- Use dplyr verbs to wrangle tsibbles