

Africast-Time Series Analysis & Forecasting Using R

<https://workshop.f4sg.org/africast/>

10. Residual diagnostics and cross validation



Outline

1 Time series cross-validation

2 Residual diagnostics

3 Recap


Outline

1 Time series cross-validation


2 Residual diagnostics

3 Recap

Issue with traditional train/test split



Time series cross-validation




Test size= forecast horizon, h

Cross-validation size=nb of experiment+ $h-1$


Time series cross-validation

Time series cross-validation



Time series cross-validation

Time series cross-validation



Creating the rolling training sets

There are three main rolling types which can be used.

- Stretch: extends a growing length window with new data.
- Slide: shifts a fixed length window through the data.
- Tile: moves a fixed length window without overlap.

Three functions to roll a tsibble: `stretch_tsibble()`, `slide_tsibble()`, and `tile_tsibble()`.

For time series cross-validation, stretching windows are most commonly used.

Time series cross-validation

Stretch with a minimum length of 24, growing by 1 each step.

```
forecast_horizon <- 12
test <- antidiabetic_drug_sale |>
  slice((n()-forecast_horizon+1):n())
train <- antidiabetic_drug_sale |>
  slice(1:(n()-forecast_horizon))
drug_sale_tcsv <- train |> slice(1:(n()-forecast_horizon)) |>
  stretch_tsibble(.init = 24, .step = 1)
```

```
# A tsibble: 2,805 x 3 [1M]
```

```
# Key:           .id [55]
```

```
  Month  Cost   .id
  <mth> <dbl> <int>
```

```
1 2000 Jan 12.5     1
```

```
2 2000 Feb  7.46     1
```

```
3 2000 Mar  8.59     1
```

```
4 2000 Apr  8.47     1
```

Time series cross-validation

Estimate RW w/ drift models for each window.

```
drug_fit_tr <- drug_sale_tcsv |>  
  model(snaive=SNAIVE(Cost))
```

```
# A mable: 55 x 2  
# Key:      .id [55]  
#   .id     snaive  
#   <int>  <model>  
1       1 <SNAIVE>  
2       2 <SNAIVE>  
3       3 <SNAIVE>  
4       4 <SNAIVE>  
# i 51 more rows
```

Time series cross-validation

Produce 8 step ahead forecasts from all models.

```
drug_fc_tr <- drug_fit_tr |>  
  forecast(h=forecast_horizon)
```

```
# A fable: 660 x 5 [1M]  
# Key:      .id, .model [55]  
   .id .model      Month  
   <int> <chr>      <mth>  
1     1  snaive 2002 Jan  
2     1  snaive 2002 Feb  
3     1  snaive 2002 Mar  
4     1  snaive 2002 Apr  
# i 656 more rows  
# i 2 more variables: Cost <dist>, .mean <dbl>
```

Time series cross-validation

```
# Cross-validated
```

```
drug_fc_tr |>
```

```
  accuracy(antidiabetic_drug_sale,  
            measures = list( point_accuracy_measures,  
                              interval_accuracy_measures,  
                              distribution_accuracy_measures))
```

```
# A tibble: 1 x 15
```

```
  .model .type    ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1  winkler  
  <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  <dbl>  
1 snaiive Test  1.48  1.90  1.56  9.09  9.59  1.17  1.20  0.254  9.73  
# i 3 more variables: scaled_pinball <dbl>, percentile <dbl>, CRPS <dbl>
```

Outline

1 Time series cross-validation

2 Residual diagnostics

3 Recap

Forecasting residuals

Residuals in forecasting: difference between observed value and its fitted value: $e_t = y_t - \hat{y}_{t|t-1}$.

Forecasting residuals

Residuals in forecasting: difference between observed value and its fitted value: $e_t = y_t - \hat{y}_{t|t-1}$.

Assumptions

- 1 $\{e_t\}$ uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2 $\{e_t\}$ have mean zero. If they don't, then forecasts are biased.

Forecasting residuals

Residuals in forecasting: difference between observed value and its fitted value: $e_t = y_t - \hat{y}_{t|t-1}$.

Assumptions

- 1 $\{e_t\}$ uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2 $\{e_t\}$ have mean zero. If they don't, then forecasts are biased.

Useful properties (for prediction intervals)

- 3 $\{e_t\}$ have constant variance.
- 4 $\{e_t\}$ are normally distributed.

Facebook closing stock price

```
fb_stock <- gafa_stock |>  
  filter(Symbol == "FB")  
fb_stock |> autoplot(Close)
```



Facebook closing stock price

```
fb_stock <- fb_stock |>
  mutate(trading_day = row_number()) |>
  update_tsibble(index = trading_day, regular = TRUE)
fit <- fb_stock |> model(NAIVE(Close))
augment(fit)
```

```
# A tsibble: 1,258 x 7 [1]
```

```
# Key:           Symbol, .model [1]
```

	Symbol	.model	trading_day	Close	.fitted	.resid	.innov
	<chr>	<chr>	<int>	<dbl>	<dbl>	<dbl>	<dbl>
1	FB	NAIVE(Close)	1	54.7	NA	NA	NA
2	FB	NAIVE(Close)	2	54.6	54.7	-0.150	-0.150
3	FB	NAIVE(Close)	3	57.2	54.6	2.64	2.64
4	FB	NAIVE(Close)	4	57.9	57.2	0.720	0.720
5	FB	NAIVE(Close)	5	58.2	57.9	0.310	0.310
6	FB	NAIVE(Close)	6	57.2	58.2	-1.01	-1.01
7	FB	NAIVE(Close)	7	57.9	57.2	0.720	0.720
8	FB	NAIVE(Close)	8	55.9	57.9	-2.03	-2.03
9	FB	NAIVE(Close)	9	57.7	55.9	1.83	1.83


Facebook closing stock price

```
augment(fit) |>  
  ggplot(aes(x = trading_day)) +  
  geom_line(aes(y = Close, colour = "Data")) +  
  geom_line(aes(y = .fitted, colour = "Fitted"))
```




Facebook closing stock price

```
augment(fit) |>  
  filter(trading_day > 1100) |>  
  ggplot(aes(x = trading_day)) +  
  geom_line(aes(y = Close, colour = "Data")) +  
  geom_line(aes(y = .fitted, colour = "Fitted"))
```




Facebook closing stock price

```
augment(fit) |>  
  autoplot(.resid) +  
  labs(x = "Day", y = "", title = "Residuals from naïve method")
```




Facebook closing stock price

```
augment(fit) |>  
  ggplot(aes(x = .resid)) +  
  geom_histogram(bins = 150) +  
  labs(title = "Histogram of residuals")
```



Facebook closing stock price

```
augment(fit) |>  
  ACF(.resid) |>  
  autoplot() + labs(title = "ACF of residuals")
```




ACF of residuals

- We assume that the residuals are white noise (uncorrelated, mean zero, constant variance). If they aren't, then there is information left in the residuals that should be used in computing forecasts.
- So a standard residual diagnostic is to check the ACF of the residuals of a forecasting method.
- We *expect* these to look like white noise.

Combined diagnostic graph

```
fit |> gg_tsresiduals()
```



Ljung-Box test


Test whether *whole set* of r_k values is significantly different from zero set.

$$Q = T(T+2) \sum_{k=1}^{\ell} (T-k)^{-1} r_k^2 \quad \text{where } \ell = \text{max lag and } T = \text{\# observations}$$

- If each r_k close to zero, Q will be **small**.
- If some r_k values large (+ or -), Q will be **large**.
- My preferences: $h = 10$ for non-seasonal data, $h = 2m$ for seasonal data.
- If data are WN and T large, $Q \sim \chi^2$ with ℓ degrees of freedom.

Ljung-Box test

$$Q = T(T + 2) \sum_{k=1}^{\ell} (T - k)^{-1} r_k^2 \quad \text{where } \ell = \text{max lag and } T = \# \text{ observations.}$$



```
# lag = h  
augment(fit) |> features(.resid, ljung_box, lag = 10)
```

```
# A tibble: 1 x 4  
  Symbol .model      lb_stat lb_pvalue  
  <chr>   <chr>      <dbl>   <dbl>  
1 FB     NAIVE(Close)  12.1     0.276
```

Outline

1 Time series cross-validation

2 Residual diagnostics

3 Recap

Recap

- 1 First, import your data and prepare them using `tsibble` function.
- 2 Visualise and see whether your series contains key patterns Use domain knowledge to understand your data and potential driving factors.
- 3 Split the data to create a training set, which you will use as an argument in your forecasting function(s). You can also create a test set to use later.
- 4 Create different rolling origins to evaluate forecast accuracy using time series cross validation

Recap

- 5 Train model to each origin
- 6 Computer forecast accuracy, use the `accuracy()` function with the `fable` as the first argument and original data as the second.
- 7 Compare methods using point, prediction interval and distributional accuracy measure; a smaller error indicates higher accuracy.
- 8 Forecast using all data for the future using the best method.
- 9 Use residual diagnostic based on residuals of the best model.