AFRICAST

TIDY TIME SERIES & FORECASTING IN R

Time series foundation models

Session 10 of 10





Outline

- Introduction
- Transfer learning
- Overview of foundation models
- Example: forecasting Australian tourism
- Evidence from public health forecasting
- Further reading and lab notebook

Foundation models

Which time series foundation models have you heard of?

Waiting for answers...



poll.nectric.com.au/s/africast-foundation-models

0 responses

● CLOSED 0 connected

SECTION 1 OF 6

Introduction

1 **Introduction** 2 Transfer learning 3 Overview of foundation models 4 Example: forecasting Australian tourism
5 Evidence from public health forecasting 6 Further reading and lab notebook



Can a pretrained model forecast our tourism series?

We want forecasts for the next four quarters of Snowy Mountains holiday travel.

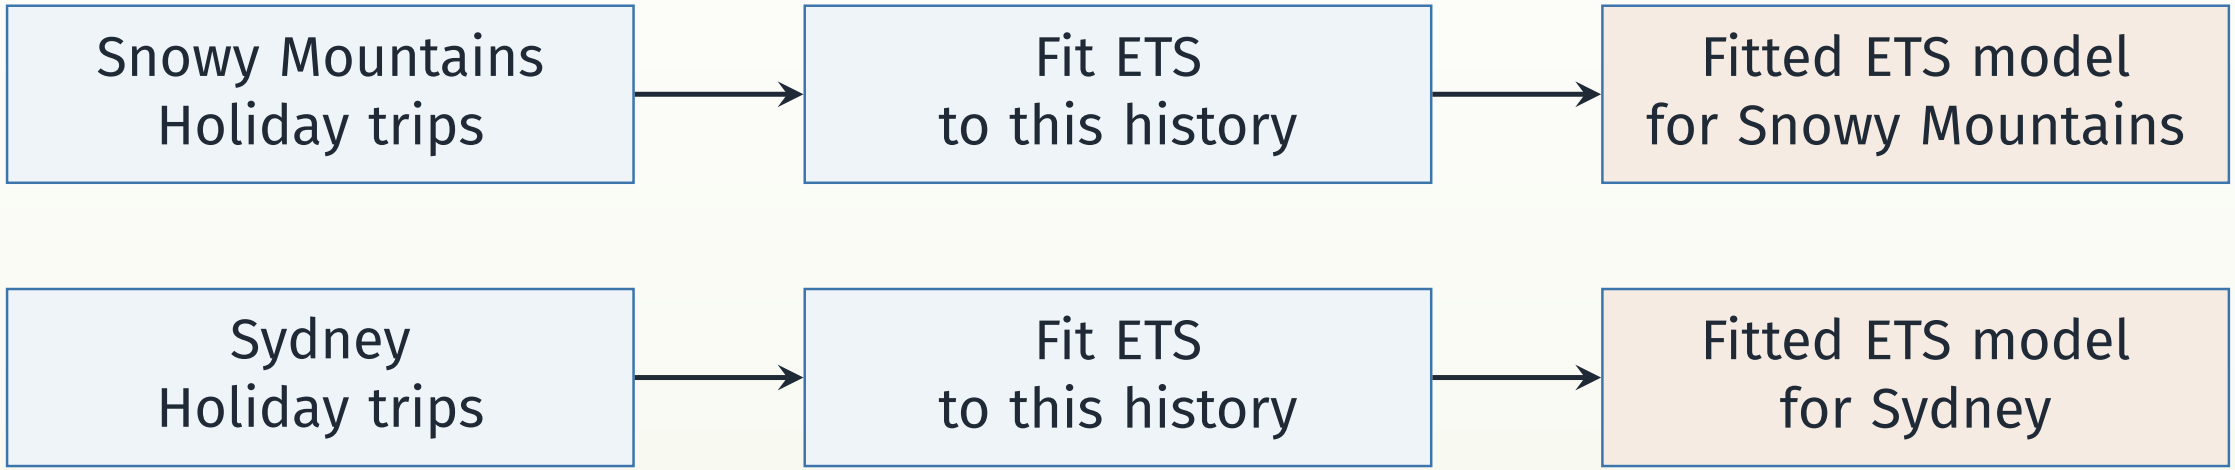
What changes in this session?

We start with a model that has already learned from many time series.

1. Supply the tourism history.
2. Examine the forecasts and their errors.
3. Try adapting the model to this history.

ETS: same method, different fitted models

We use the same fitting method, but learn from each region's history separately.



What changes with a foundation model?

We can reuse a model already trained on many series, then supply our tourism history to obtain a forecast.

What is a foundation model?

Train broadly, then reuse

A foundation model is trained on large, varied data so that it can be reused for a range of tasks and settings.

- The initial training is called **pretraining**.
- For forecasting, this training uses a collection of time series.
- We supply our tourism history and ask the trained model for the next four quarters.

Four characteristics to remember

1. Large, varied training data
Learn from many examples and settings.

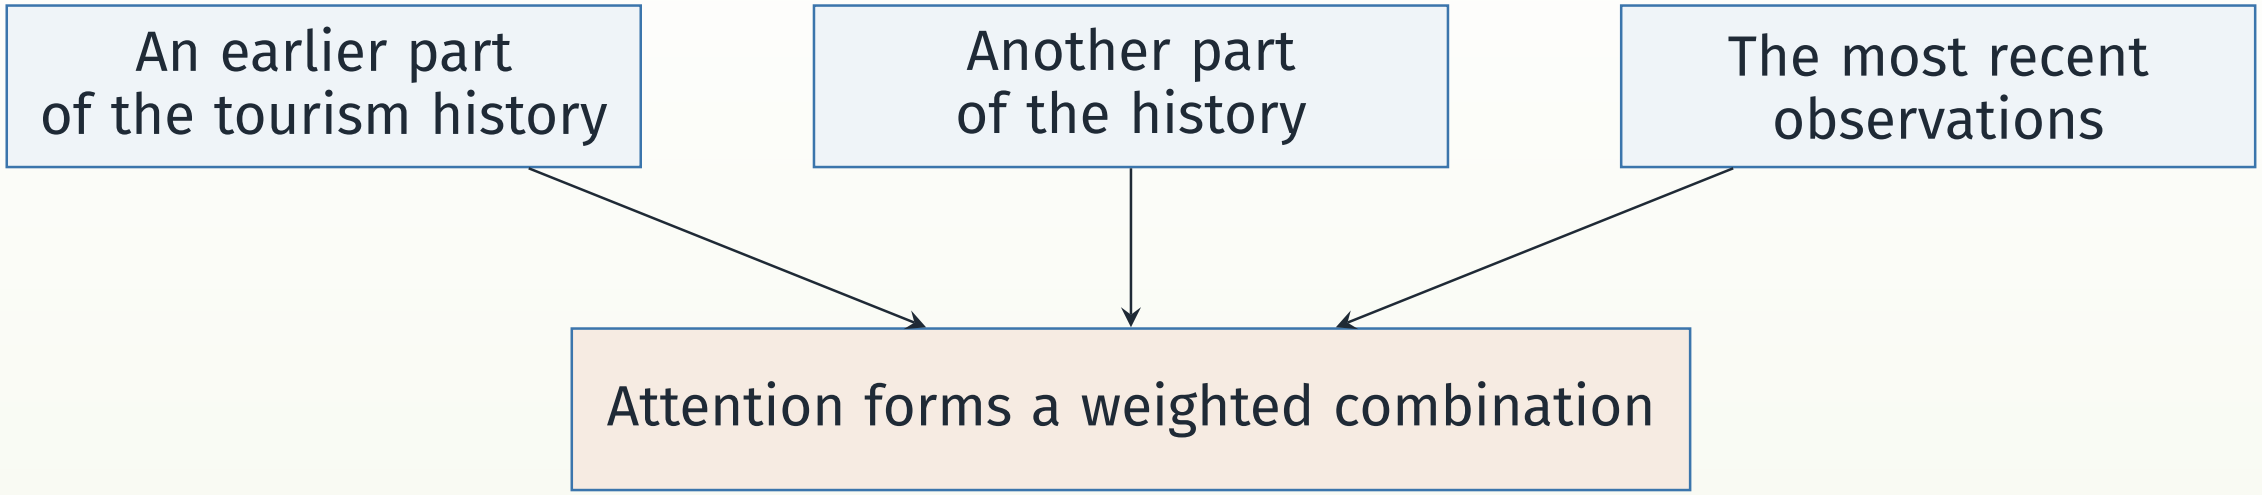
2. Often many parameters
Learned numbers store the model's behaviour.

3. Reusable across tasks
Apply the trained model to different series.

4. Can be adapted
Further training can tailor it to our data.

These features make reuse possible. They do not guarantee the best forecast for every series.

Attention combines information from the history



- Different parts of the history can receive different weights.
- A **transformer** uses attention and neural-network layers to process this information.

Two common routes to foundation-model forecasting

1. Models built for time series

Learn from collections of time series, then forecast a new series.

TimeGPT, Chronos-2, TimesFM and Moirai follow this route.

2. Language models adapted for forecasting

Start with a pretrained language model and adapt how time-series inputs and forecasts are represented.

Examples: PromptCast and Time-LLM. Our code uses the first route.

Advantages and disadvantages

Advantages	Disadvantages
Use out-of-the-box forecasting pipelines	Privacy concerns
Can generate forecasts with few data points	No control of model's capabilities
Require less forecasting expertise	May not be optimal for every use case
Reusable across tasks	High resource requirements

For our tourism example

Getting a forecast is easier. We still need to check whether it is useful.

SECTION 2 OF 6

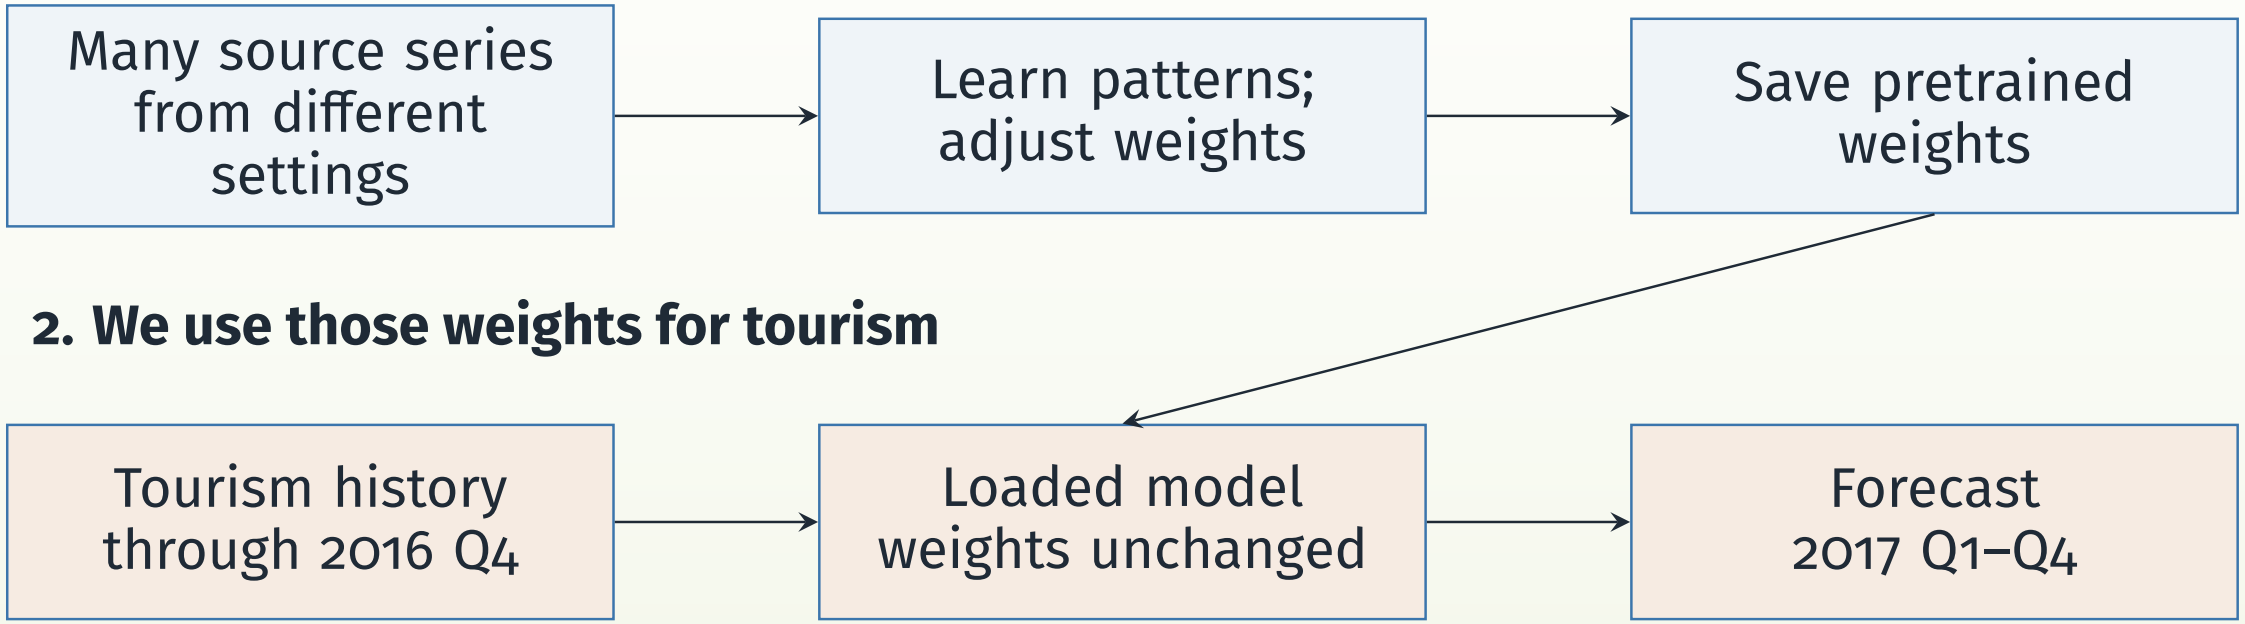
Transfer learning

1 Introduction
2 **Transfer learning**
3 Overview of foundation models
4 Example: forecasting Australian tourism
5 Evidence from public health forecasting
6 Further reading and lab notebook



Pretraining and forecasting are separate stages

1. Developers train the model



2. We use those weights for tourism

Using a model learned from other series is **transfer learning**.

What makes building a foundation model difficult?

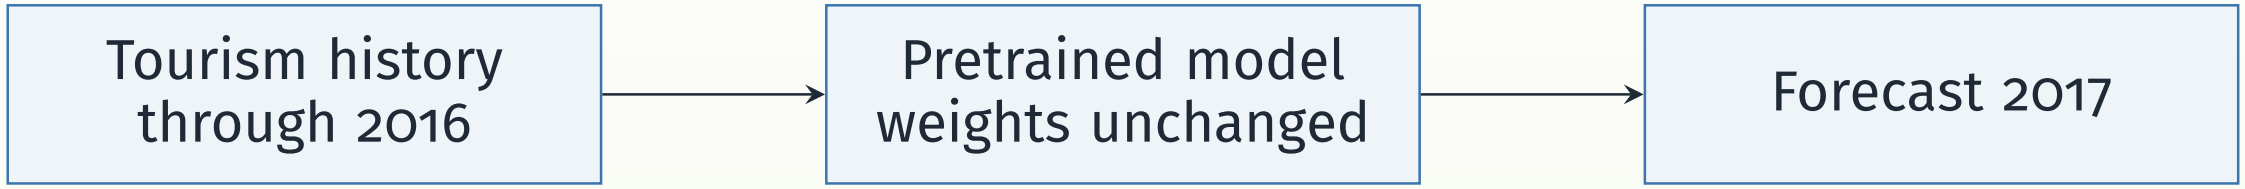
Challenge	What the model developer must handle
Different frequencies and domains	Quarterly tourism and hourly energy demand have different patterns.
Different forecast horizons	One week means 7 daily values or 168 hourly values.
Training resources	Large-scale training takes time, memory and often GPUs.
Enough varied training data	Open time-series data can be scarce; combining sources takes work.

Why start with pretrained weights?

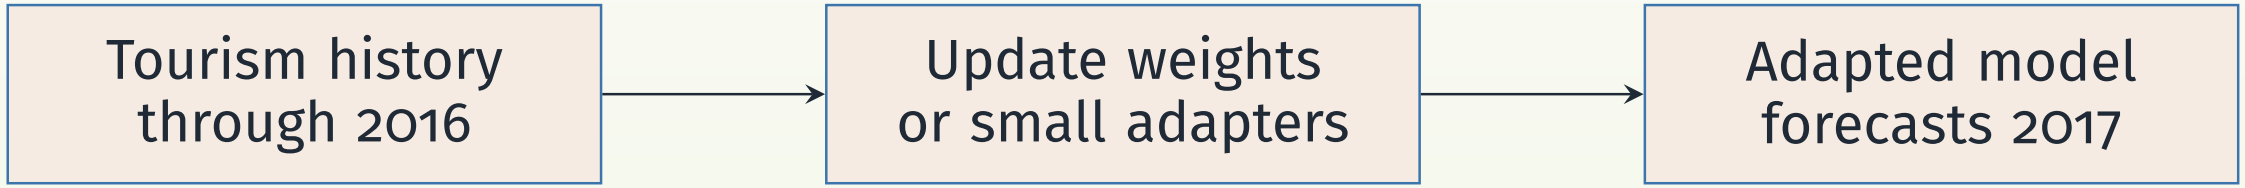
We can reuse that training effort, then test the model on our own task.

Zero-shot or fine-tuning: what changes?

Zero-shot: use the pretrained weights



Fine-tuning: start from those weights and adapt



Both use tourism history. Neither is allowed to learn from the 2017 test observations.

SECTION 3 OF 6

Overview of foundation models

1 Introduction
2 Transfer learning
3 **Overview of foundation models**
4 Example: forecasting Australian tourism
5 Evidence from public health forecasting
6 Further reading and lab notebook

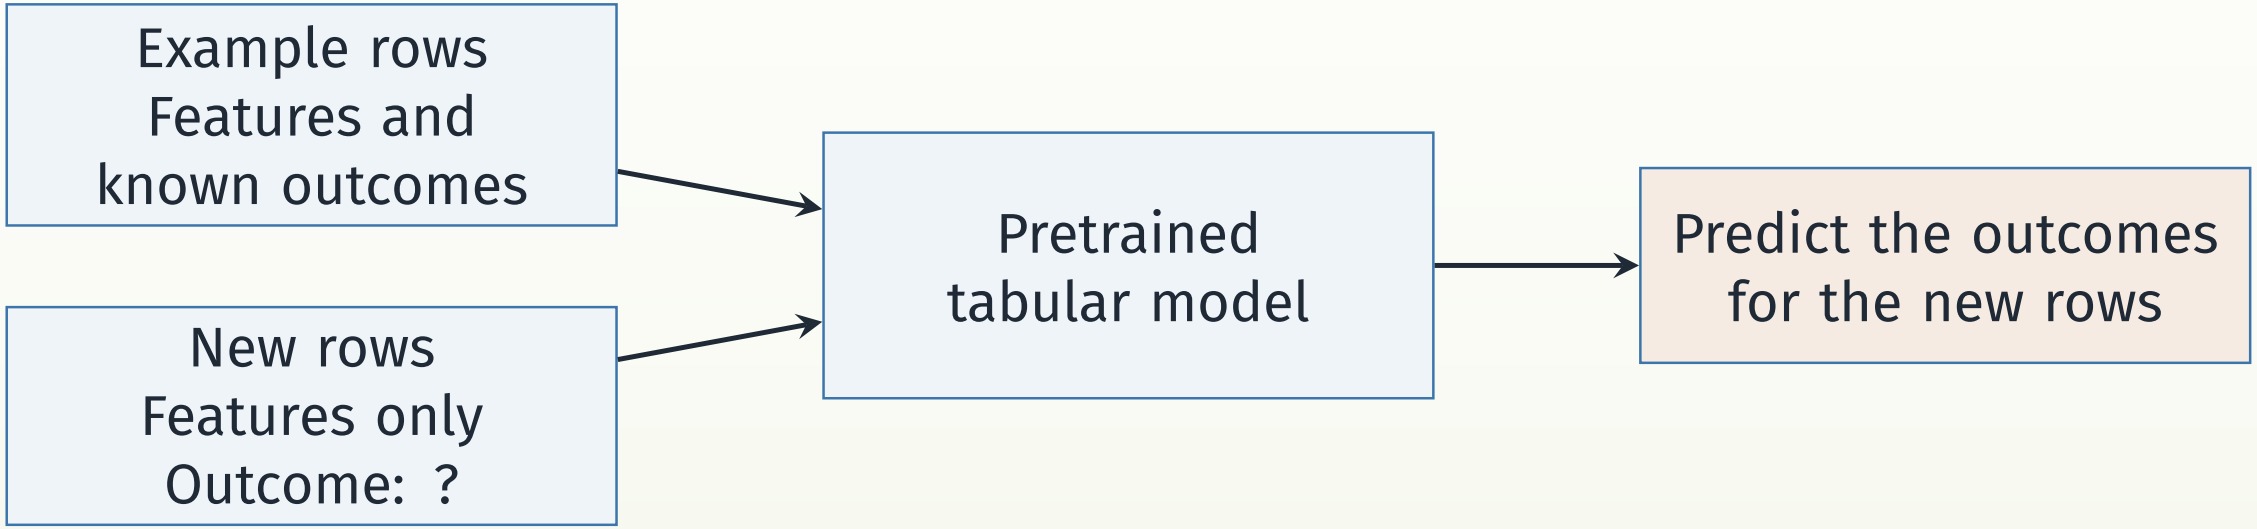


Comparing foundation forecasting models

Model	Architecture	Input representation	Forecast output
TimeGPT-1	Encoder–decoder	History and positions	Point and intervals (API)
Chronos-2	Encoder-only	Patches; optional predictors	Quantiles
TimesFM 2.5	Decoder-only	Historical patches	Point and quantiles
Moirai 1.0	Encoder-only	Patches; optional predictors	Distribution
Lag-Llama	Decoder-only	Lags and time features	Distribution
Tiny Time Mixers	MLP mixer	Patches; predictors when adapted	Point forecasts

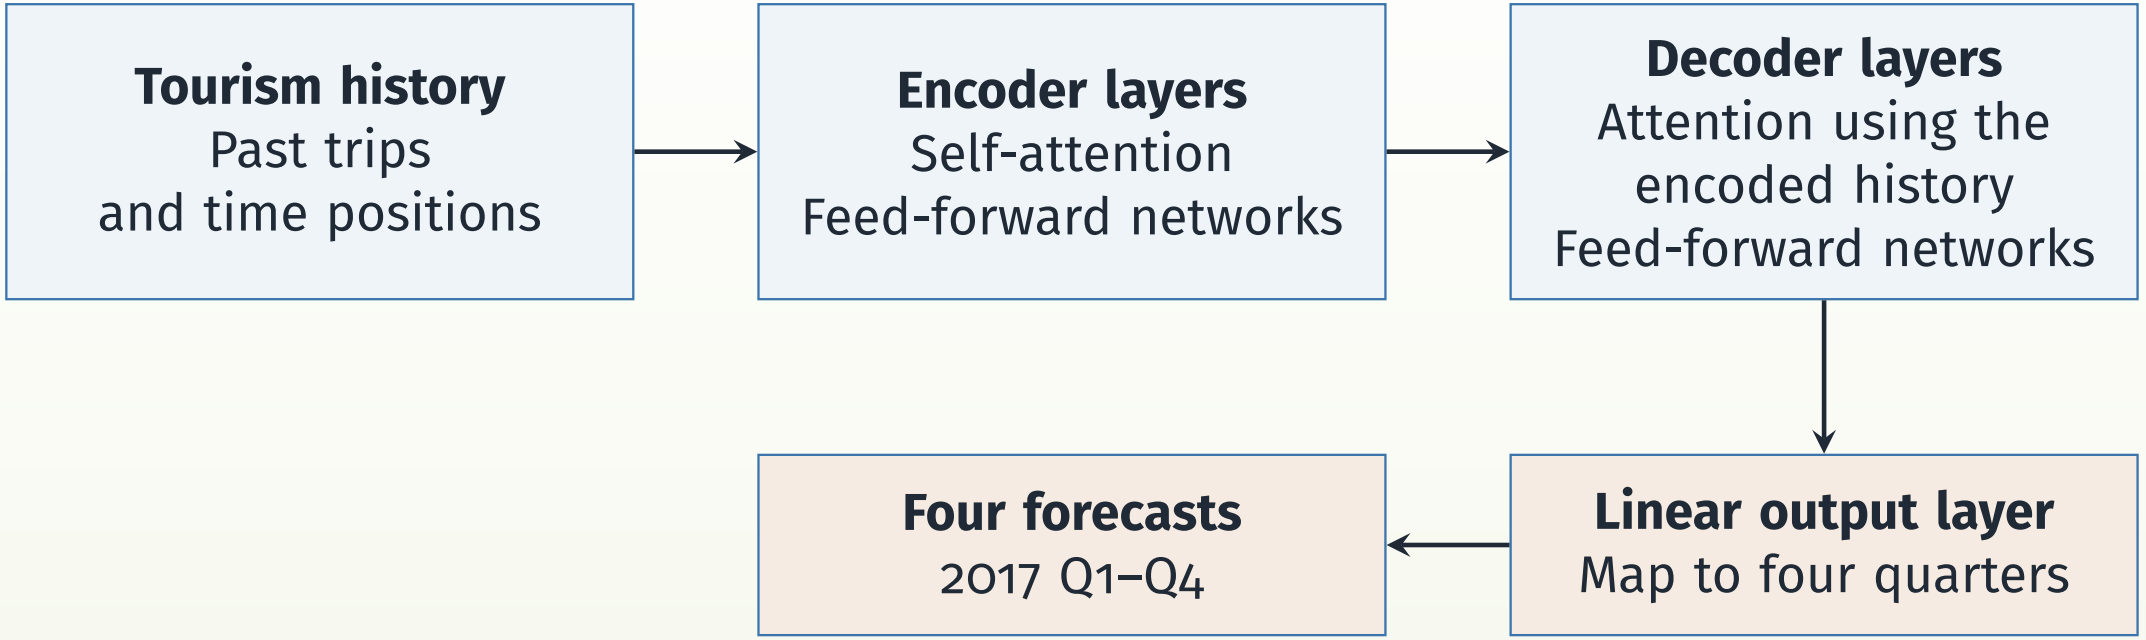
Foundation models can also work with tables

Examples: **TabPFN** and **TabICL**.



- Pretrain on many tasks, often using simulated tables.
- Use example rows as **context**; keep pretrained weights fixed.
- For forecasting, turn time and calendar information into features.

TimeGPT: an encoder–decoder transformer



Attention combines information across positions. A **feed-forward network** refines the information at each position.

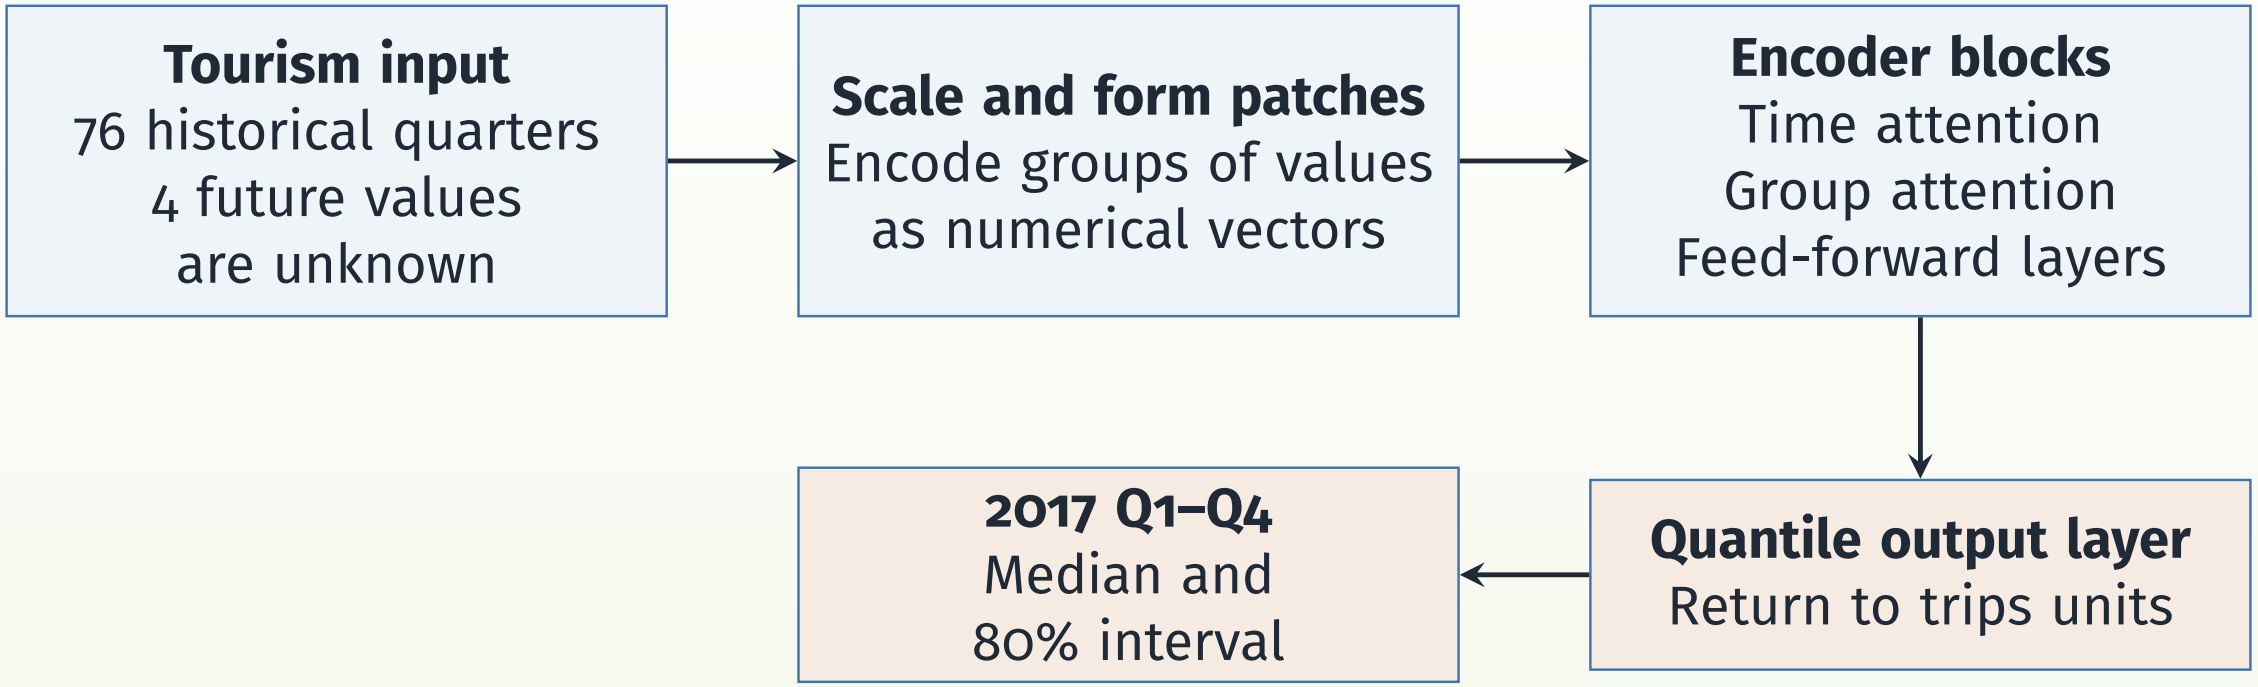
Simplified TimeGPT-1 architecture (Garza et al.). Repeated layers and normalisation are omitted.

TimeGPT: the model runs on Nixtla’s service



- For tourism, send the history and ask for four quarters.
- Set fine-tuning arguments to adapt the model to that history.
- The call needs an API key and available account quota.

Chronos-2: an encoder-only transformer



Time attention: share information across the history.
Group attention: share information across related series or predictors, when supplied.

10-minute break

After the break

Forecast quarterly tourism, examine outputs, and compare zero-shot forecasts with fine-tuned forecasts.

SECTION 4 OF 6

Example: forecasting Australian tourism

1 Introduction
2 Transfer learning
3 Overview of foundation models
4 **Example: forecasting Australian tourism**
5 Evidence from public health forecasting
6 Further reading and lab notebook



Our forecasting task

- **Series:** Snowy Mountains, Holiday travel.
- **Target:** quarterly overnight trips, in thousands.
- **History:** 1998 Q1 to 2017 Q4, 80 observations.
- **Task:** forecast the four quarters of 2017 using data available through 2016.

The steps we will follow

Prepare the history → forecast → evaluate → adapt → compare.

From a tsibble to model input

Select the same series in R. We will pass it to the models through their Python interfaces:

```
tourism |>
  filter(Region == "Snowy Mountains", Purpose == "Holiday")
```

The Python interface uses three columns:

Workshop column	Model column	Meaning
Region and Purpose	`unique_id`	Series identifier
Quarter	`ds`	Quarter-start date
Trips	`y`	Trips, thousands

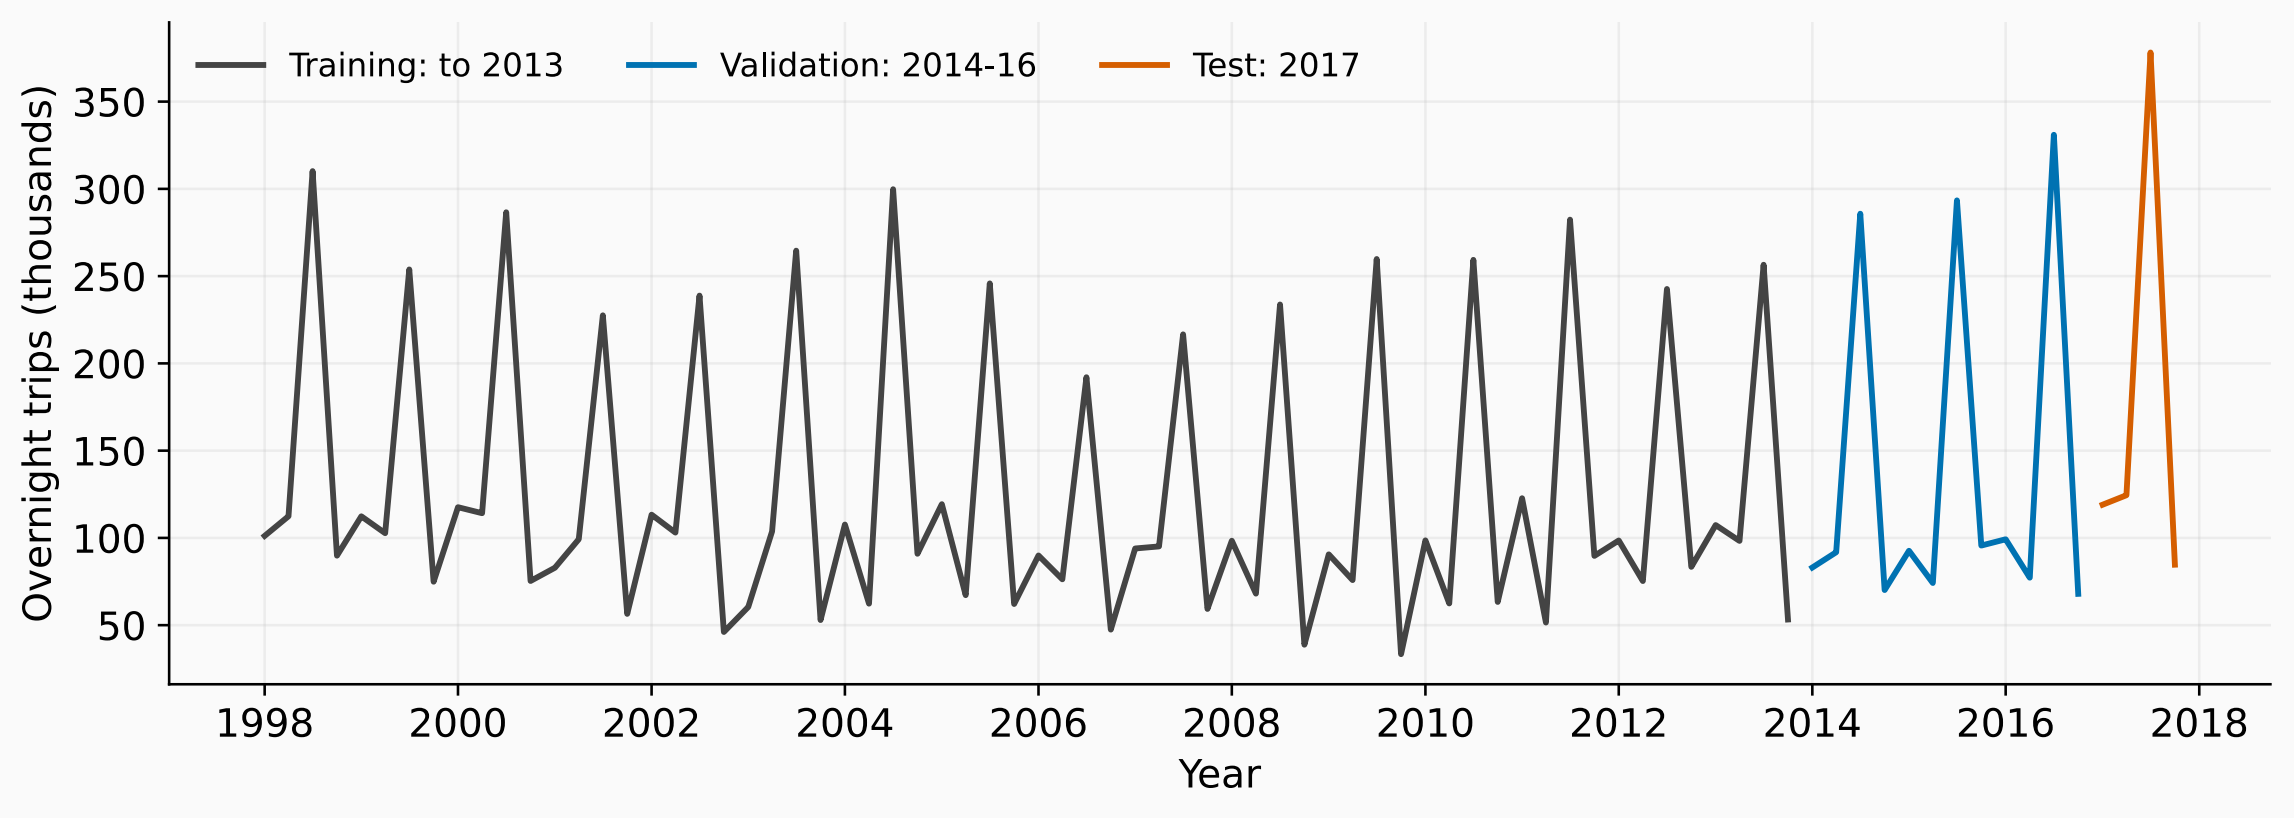
Prepare the tourism series in Python

The model needs an identifier, a date and the number of trips.

```
import pandas as pd
import numpy as np
tourism = pd.read_csv("data/tourism.csv", parse_dates=["Quarter"])
keep = ((tourism.Region == "Snowy Mountains") &
        (tourism.Purpose == "Holiday"))
df = tourism.loc[keep].rename(columns={"Quarter": "ds", "Trips": "y"})
df["unique_id"] = "Snowy Mountains / Holiday"
df = df[["unique_id", "ds", "y"]].sort_values("ds")
H = 4
history, test = df.iloc[:-H], df.iloc[-H:]
```

Result: 76 historical observations and 4 test observations.

History, validation and the final test



Validation uses three four-quarter forecasts covering 2014–2016. The 2017 observations are reserved for the final comparison.

Start with a seasonal-naive forecast

For each quarter of 2017, repeat the corresponding observation from 2016.

```
baseline = test[["unique_id", "ds"]].copy()
baseline["prediction"] = history["y"].tail(H).to_numpy()
baseline["prediction"].round(2).tolist()
```

```
[99.21, 77.21, 330.99, 67.95]
```

These are forecasts for Q1–Q4, in thousands of trips. Every model will be compared with this benchmark on the same dates.

Zero-shot forecasting with TimeGPT

```
from nixtla import NixtlaClient

client = NixtlaClient()  # reads NIXTLA_API_KEY
timegpt = client.forecast(
    df=history, h=H, freq="QS", level=[80],
    finetune_steps=0,  # no update to model weights
)
```

`h=4`: four quarters; `QS`: quarter-start dates; `level=[80]`: an 80% interval.

Zero-shot forecasting with Chronos-2

Load the model and request a median forecast with an 80% interval.

```
from chronos import Chronos2Pipeline

pipeline = Chronos2Pipeline.from_pretrained(
    "amazon/chronos-2", device_map="cpu"
)
chronos = pipeline.predict_df(
    history, prediction_length=H,
    quantile_levels=[0.1, 0.5, 0.9],
    id_column="unique_id", timestamp_column="ds",
    target="y",
)
```

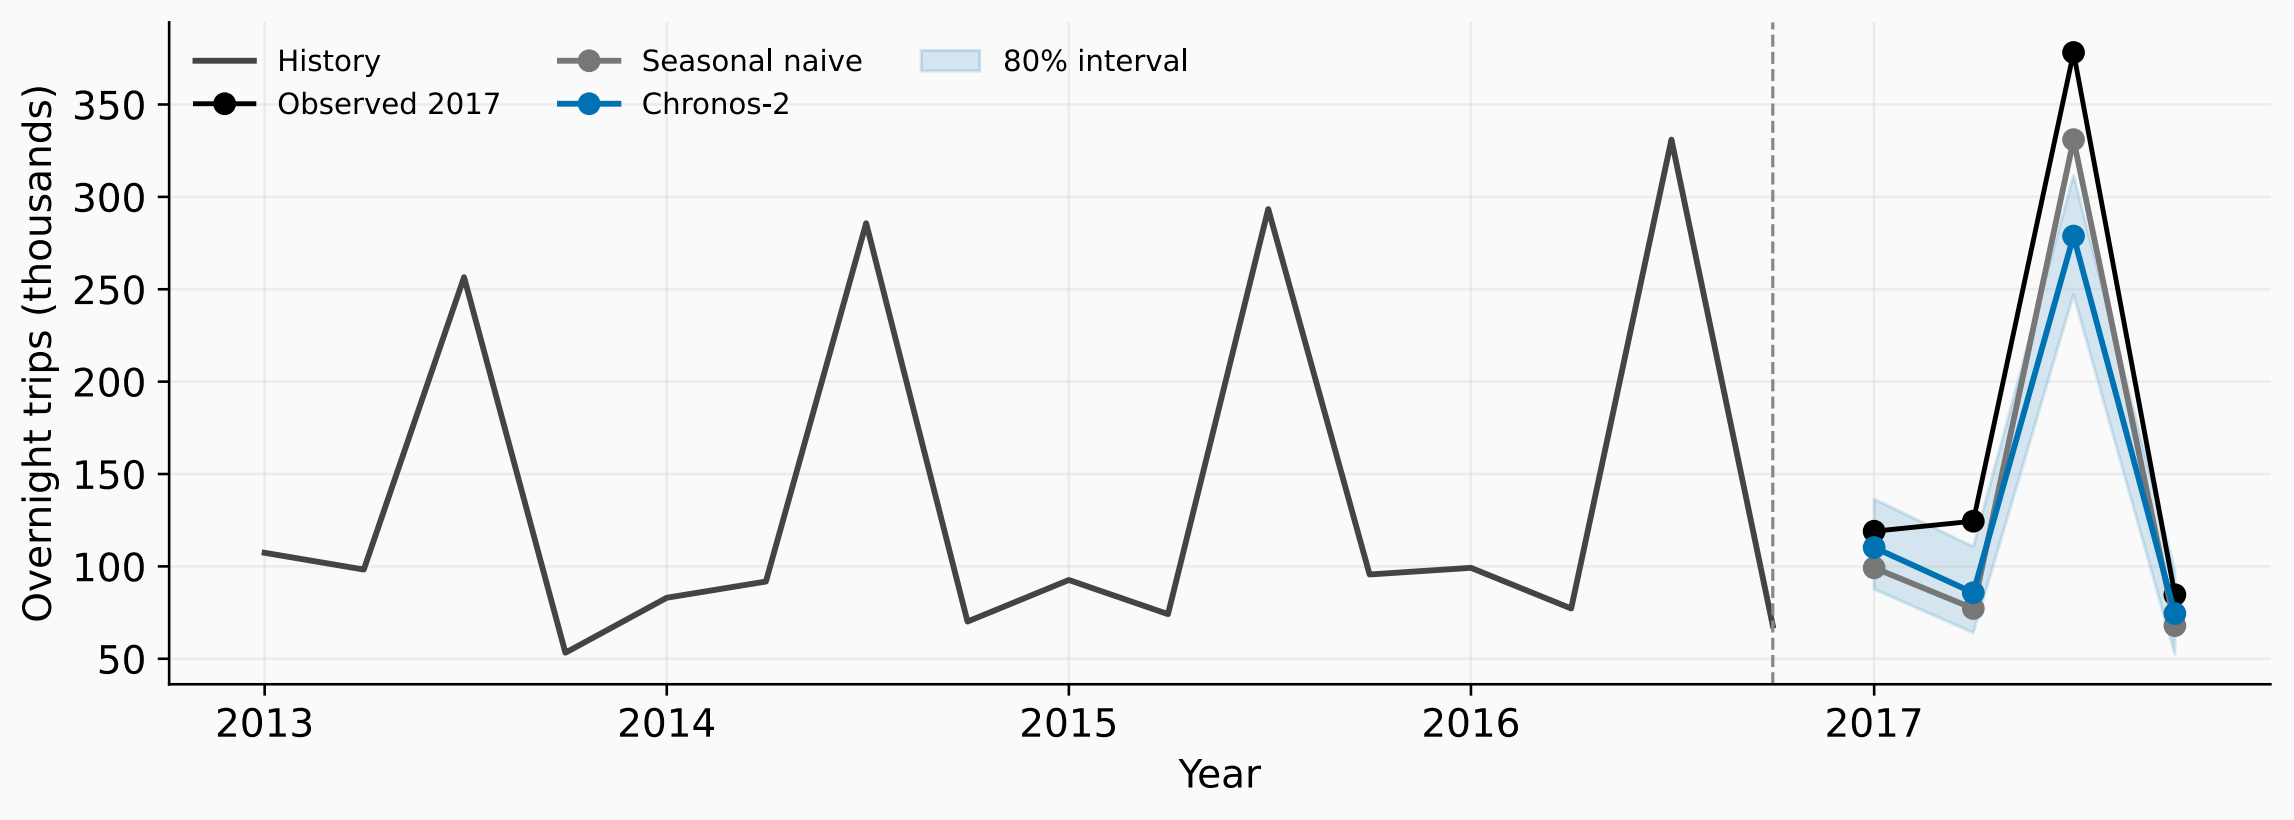
What does the Chronos-2 output contain?

```
chronos[["ds", "0.1", "0.5", "0.9"]].round(2)
```

ds	0.1	0.5	0.9
2017-01-01	87.67	110.23	136.42
2017-04-01	64.07	85.74	110.65
2017-07-01	247.12	278.78	311.71
2017-10-01	52.49	74.47	97.20

`0.5` is the median forecast. `0.1` and `0.9` are the lower and upper interval limits. Values are thousands of trips.

Chronos-2: forecasts and intervals



The model misses the winter peak. Its 80% interval contains 2 of the 4 observed quarters; four quarters are too few to assess long-run coverage.

MAE or RMSE?

Let e be the observed number of trips minus the forecast.

Measure	Calculation	How to read it
MAE	$\text{mean}(\lvert e \rvert)$	Average absolute miss
RMSE	$\sqrt{\text{mean}(e^2)}$	Gives large misses more weight

- Both are in thousands of trips; smaller is better.
- Keep MAE as the main measure for our median forecasts.
- Also report RMSE to show the effect of large errors.

Calculate both errors on the same forecasts

Match each forecast with the observation for that quarter.

```
checked = test.merge(chronos, on=["unique_id", "ds"],
                     validate="one_to_one")
errors = checked["y"] - checked["0.5"]
mae = errors.abs().mean()
rmse = np.sqrt((errors ** 2).mean())
print(f"MAE: {mae:.2f}; RMSE: {rmse:.2f}")
```

```
MAE: 39.27; RMSE: 53.75
```

Before changing the model, check its performance on earlier forecast periods as well.

Check earlier forecasts before choosing settings

Forecast one year at a time, using only the history available then.

History ends	Forecast year
2013 Q4	2014
2014 Q4	2015
2015 Q4	2016

Model	Validation MAE	Validation RMSE
Seasonal naive	17.68	20.65
Chronos-2	18.45	23.13

These scores cover 12 validation quarters. Keep 2017 for the final test.

Fine-tuning TimeGPT

Use the same history and horizon, with ten training updates.

```
timegpt_ft = client.forecast(
    df=history, h=H, freq="QS", level=[80],
    finetune_steps=10,
    finetune_depth=1,
    finetune_loss="mae",
)
```

- `finetune_steps` counts training updates; `h` counts future quarters.
- `finetune_loss="mae"` uses absolute errors during adaptation.
- Compare on the same four quarters using both metrics.

Fine-tuning Chronos-2 with LoRA

LoRA trains a small set of added weights while keeping the original weights fixed.

```
adapted = pipeline.fit(
    inputs=[history.y.to_numpy(dtype=np.float32)],
    prediction_length=H,
    finetune_mode="lora", num_steps=10,
    **lora_settings,
)
```

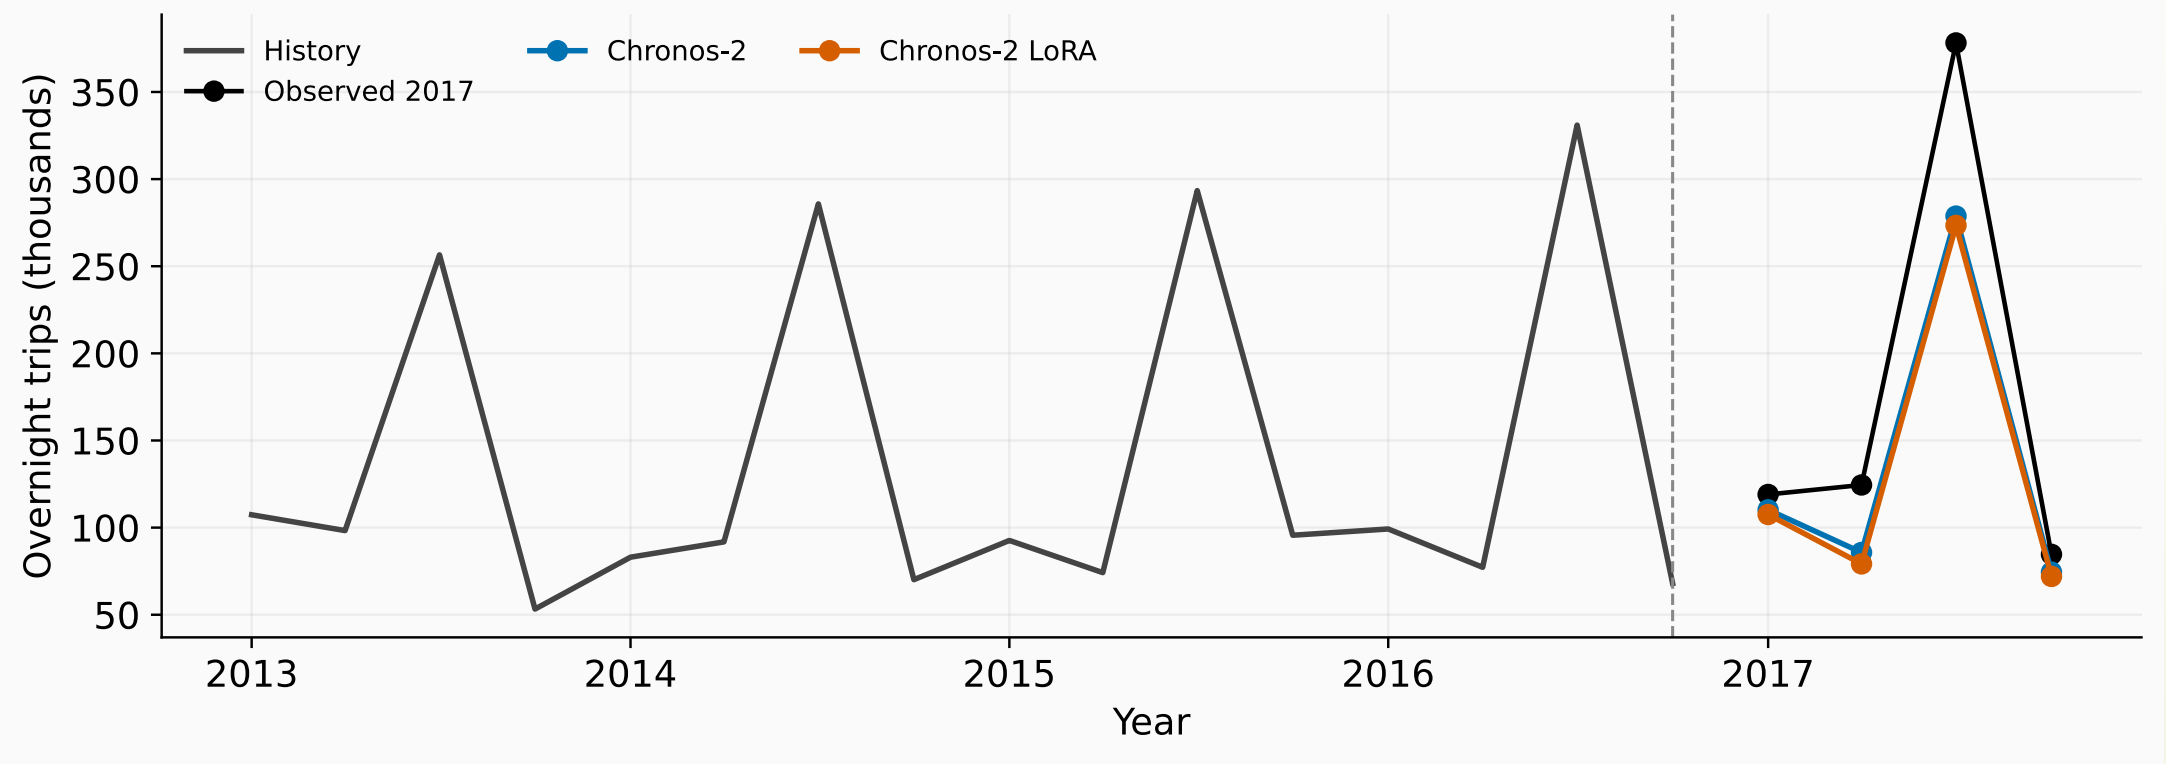
`lora_settings` contains the fixed training settings listed in the notebook.

What changes?

We get an adapted model. The history and the four forecast dates stay the same.

Forecast again with the adapted model

```
chronos_ft = adapted.predict_df(
    history, prediction_length=H, quantile_levels=[0.1, 0.5, 0.9],
    id_column="unique_id", timestamp_column="ds", target="y",
)
```



The adapted model still misses the winter peak. Did its overall error improve?

Results on the held-out tourism year

Model	MAE	RMSE
Seasonal naive	32.73	35.80
Chronos-2	39.27	53.75
Chronos-2 LoRA	43.55	57.68

Both measures: thousands of trips. Chronos-2 rows use median forecasts.

Seasonal naive has the smallest error under both measures. Ten-step LoRA made both scores worse; this one year cannot establish a general ranking.

What other information is known when we forecast?

Suppose we are forecasting in December 2016.

Information for 2017	Available now?
Calendar quarter	Yes
Actual weather	No
Actual number of trips	No: our forecast target

Supply historical predictor values and any future values known at forecast time.

The lab uses calendar features through `future_df`. Check their value on earlier forecasts.

SECTION 5 OF 6

Evidence from public health forecasting

1 Introduction
2 Transfer learning
3 Overview of foundation models
4 Example: forecasting Australian tourism
5 **Evidence from public health forecasting**
6 Further reading and lab notebook



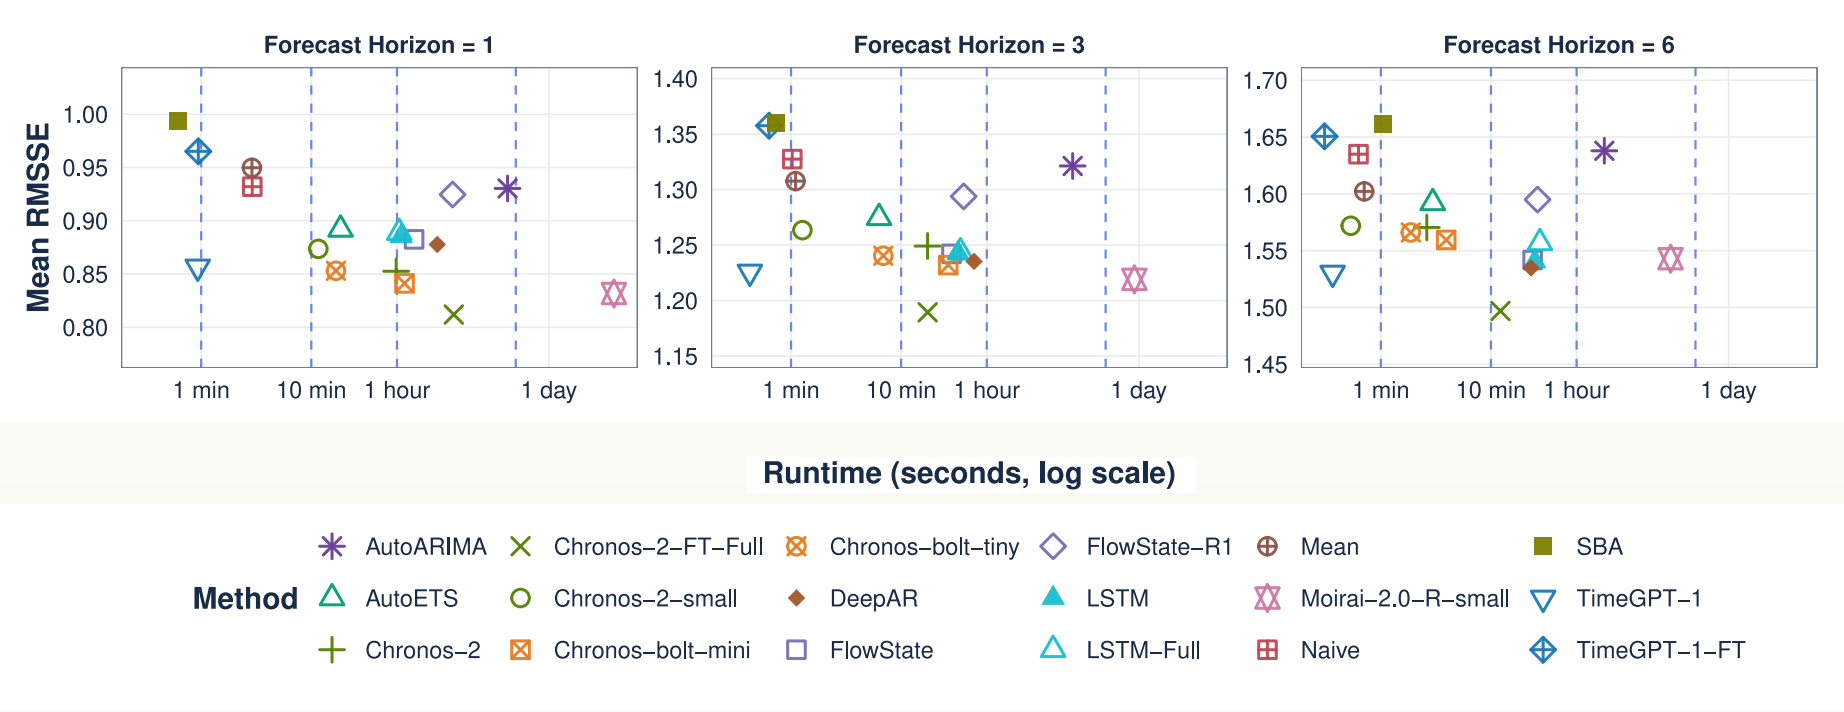
Beyond tourism: forecasts for health supplies

- 20,736 monthly product-location series from Lao.
- Five clinical domains; forecast horizons of 1, 3 and 6 months.
- Compare pretrained and fine-tuned models with statistical and global forecasting benchmarks.

Why use RMSSE here?

RMSSE is RMSE divided by a scale based on historical seasonal changes. It lets us compare errors across series with different volumes. Smaller is better.

Lao health supplies: accuracy and runtime



Lower and further left is preferable. Chronos-2 fine-tuning improves accuracy here, but adds computation.

Source: Hewage, Rostami-Tabar, Syntetos and Liberatore, *Foundation models in low-resource settings*.

SECTION 6 OF 6

Further reading and lab notebook

1 Introduction
2 Transfer learning
3 Overview of foundation models
4 Example: forecasting Australian tourism
5 Evidence from public health forecasting
6 **Further reading and lab notebook**



Continue with the tourism lab

- Read and reproduce saved outputs first.
- Run models when credentials or weights are available.
- Try another region and purpose, using earlier validation periods.
- Record the version, data split, metrics and one limitation.

Independent lab

`exercises/day5-foundation_models/tourism_lab.ipynb`

Further reading

- FPP Python, Chapter 15: foundation models and transfer learning.
- Peixeiro, *Time Series Forecasting Using Foundation Models*, Chapters 1 and 2: concepts and the pretraining-to-adaptation flow.
- Nixtla and Chronos: executable APIs.
- Bergmeir, *Forecasting for Data Scientists*, Theory Session 2: global models and evaluation.
- Molnar (2026), *Tabular Foundation Models*, Chapters 3–4 and 9: in-context prediction and forecasting.
- Tourism documentation. We use the workshop's CSV snapshot.
- Harsha Chamara Hewage, Bahman Rostami-Tabar, Aris Syntetos and Federico Liberatore. *Foundation models in low-resource settings: Evaluating AI-driven forecasting in global health supply chains.*

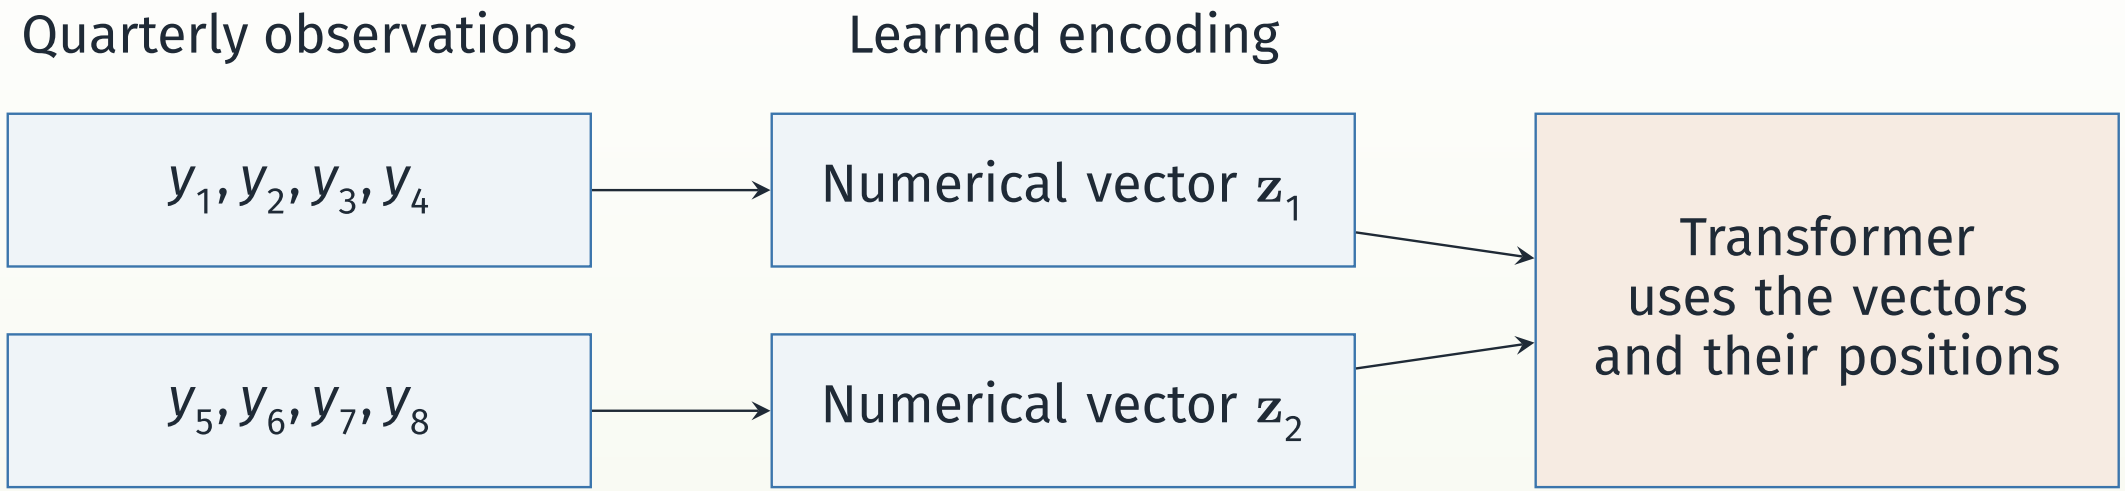
Optional: turn tourism forecasting into a table

One row per quarter; predict trips from the feature columns.

Quarter	Time index	Quarter of year	Trips (thousands)
2016 Q1	73	1	99.21
2016 Q2	74	2	77.21
2016 Q3	75	3	330.99
2016 Q4	76	4	67.95
2017 Q1	77	1	?
2017 Q2	78	2	?
2017 Q3	79	3	?
2017 Q4	80	4	?

Use historical rows as context. Supply future features; keep future trips unknown.

Optional: patches and embeddings

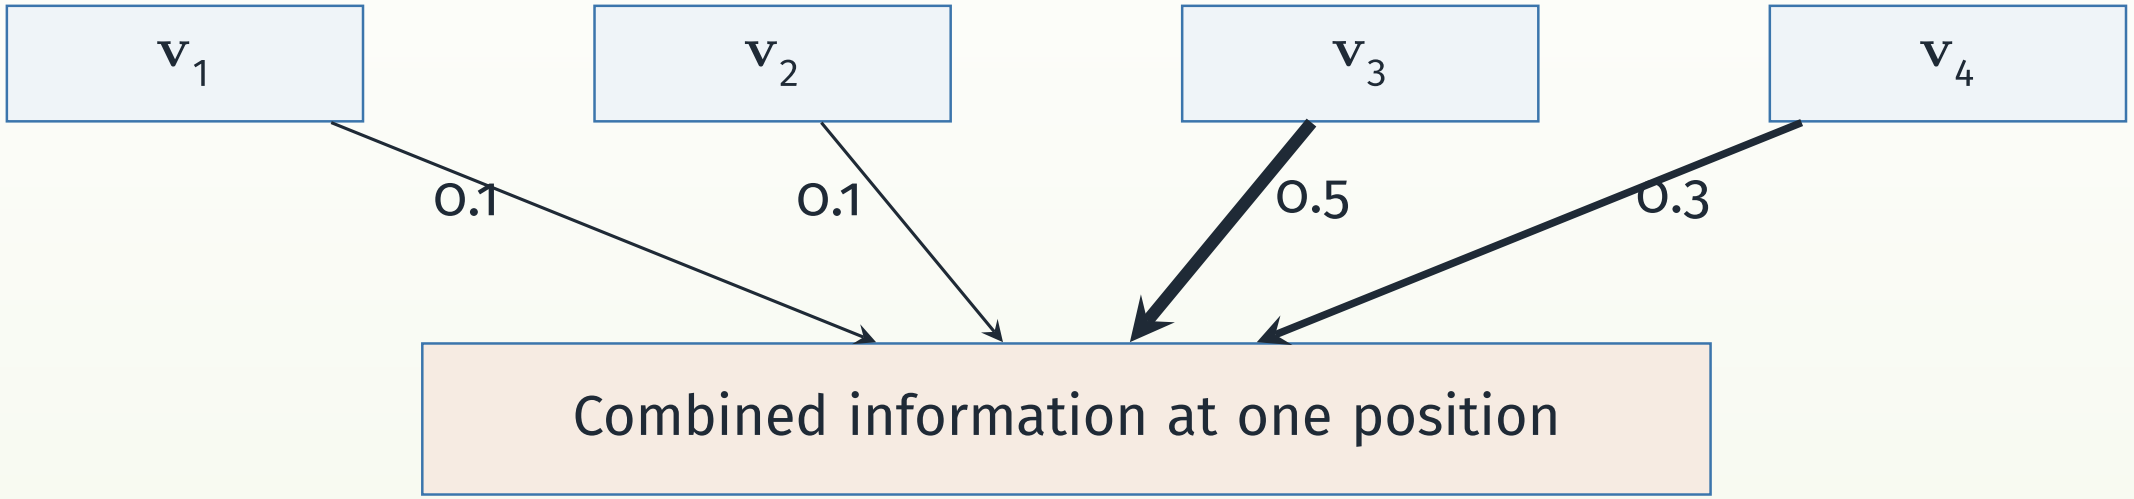


- A **patch** is a group of consecutive observations.
- An **embedding** turns that patch into a list of numbers the model can process.
- Position information preserves the order of the patches.

Four-quarter patches are illustrative. Model patch sizes vary.

Optional: attention weights

Attention combines information from past patches.

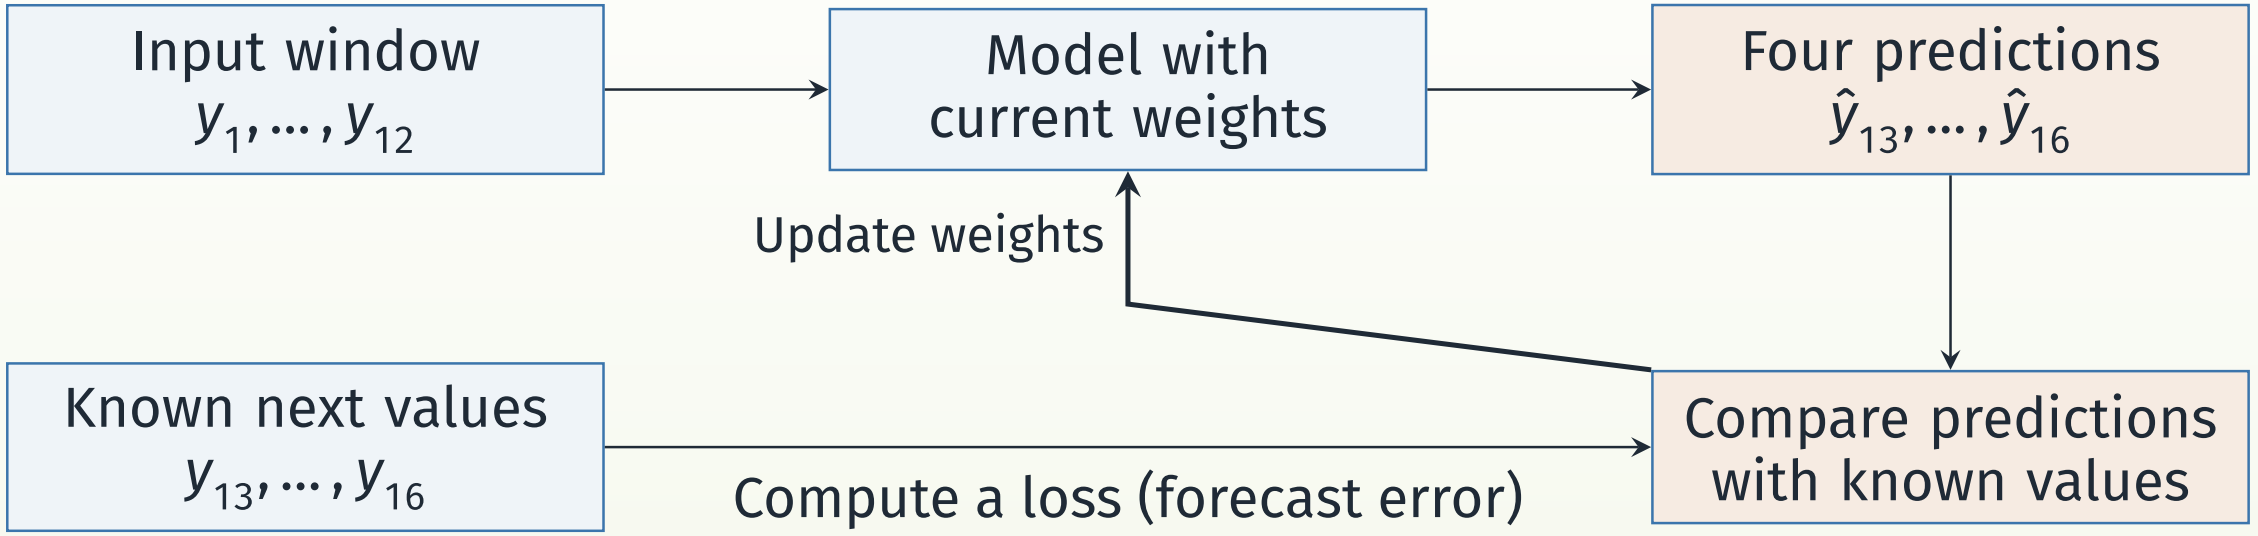


- Here the combination is $0.1\mathbf{v}_1 + 0.1\mathbf{v}_2 + 0.5\mathbf{v}_3 + 0.3\mathbf{v}_4$.
- Each attention head forms its own combination.

These illustrative weights depend on the input. They are not stored model parameters.

Optional: historical training windows

Use a past window to predict the observations immediately after it.



Repeat across windows and series. The training targets are already in the historical data.